

Computational Self-Awareness in Musical Robotic Systems

*Endowing musical robots with
self-awareness — an experimental and
exploratory study*

David Thorvaldsen



Thesis submitted for the degree of
Master in Informatics: Robotics and Intelligent
Systems
60 credits

Institute for Informatics
Faculty of mathematics and natural sciences

UNIVERSITY OF OSLO

Spring 2022

Computational Self-Awareness in Musical Robotic Systems

*Endowing musical robots with
self-awareness — an experimental and
exploratory study*

David Thorvaldsen

© 2022 David Thorvaldsen

Computational Self-Awareness in Musical Robotic Systems

<http://www.duo.uio.no/>

Printed: Reprocentralen, University of Oslo

Abstract

Engineering a computing system for a certain environment often requires some knowledge of said environment; both on the end of the system-creator, as well as for the computing system in turn. This knowledge will then have to be used in the environment the computing system is situated in, in all the situations and possible states the system can be in during its lifetime. However, foreseeing and predicting, at design-time, all possible future states and scenarios a computing system will be in during its lifetime is often hard, and sometimes impossible (in the case of coincidental faults e.g.). Increasingly extreme, complex, dynamic, and ever-changing environments, simply enlargens this challenge. This calls for online and continuous learning at run time, where computing systems are themselves able to observe, learn, adapt, and act on their own, independently from their creator.

Especially within collectives of multiple robots [17, 27], communication between individuals—even though they might be autonomous individual robots—is still a challenge one has to deal with when designing them. Coordination, for humans and robots alike, often tend to demand a lot of communication. To add on top of this, when networks or collectives of sub-systems like wireless sensors are communicating with each other, they often have their own internal clock which often do not align. Hence, we see that both the magnitude of communication needed, as well as the inherent challenges with communication, call for effective approaches to communication with coordination as the goal.

As K. Konishi & H. Kokame [11] points out, one technical problem in e.g. wireless sensor networks is that each sensor node rely on accurate internal and synchronized clocks, especially from the perspective of sensor fusion and co-ordinating communication among nodes. Thus, time synchronization becomes recognized as one of the crucial problems in distributed and wireless systems [28]. Various attempts at synchronizing messages and communication, in order to e.g. align messages exchanged at unsynchronized intervals, have been made [28]; however, such attempts and protocols often require the computation of message exchange and processing, which wastes the limited computation capability of nodes and causes communication delays. To this, Konishi & Kokame [11] also point out how for synchronized pulse coupled oscillators (PCOs), such computation is not required. If internal clocks in communicating nodes are instead altered and synchronized to other nodes's clocks—instead of letting all clocks stay unsynchronized and constantly trying to compensate for this—it is rather apparent that getting rid of this need for compensation will considerably reduce computation needs while synchronized communication is still achieved. Hence, by achieving time synchronization, both the computational load, in addition to the need for communication, is considerably reduced.

Thus, by achieving time synchronization [28] through synchronizing pulse coupled oscillators, both the computational load, in addition to the need for communication, is considerably reduced; hence freeing up and saving valuable resources like e.g. processing power and energy (battery) in networks of individuals and possibly autonomous nodes, agents, or robots.

Despite its elusive, and at times hard-to-tangibly-define, nature; the psychological concept of Self-Awareness has relatively recently served as a rich source of new inspiration and conceptual tools and frameworks, previously unknown or undiscovered by the computer engineer.

With self awareness, online and continuous learning is achieved to a higher degree in contrast with other approaches (like ODA and MAPE-K), due to the limitations and downsides of these older approaches, as well as the advantages and upsides to considering computational self-awareness in computing systems. If one wants to achieve continuous adaptation of a system or of system-components (e.g. in a collective) — some sort of intelligence might be necessary to endow it with. Endowing computing systems with Self-Awareness can be beneficial in several respects, including but not limited to a greater capacity to adapt, to build potential for future adaptation in unknown environments, and to explain their behaviour to humans and other systems [13].

Taking inspiration from the fascinating natural phenomena of self synchronizing fireflies, a synchronization simulator imitating and modelling this process for—instead of fireflies—a collective of musical robots, has been designed and tested. It has been found that e.g. the sensitivity of the robots when it comes to adjusting themselves on the way to synchronizing to each other does not have as much a say on the stability nor performance of the synchronization task as e.g. the robot collective size does. Furthermore, it is also found that the synchronization task can still be achieved despite more relaxed connectivity in the network of robots.

Acknowledgements

Tusen takk til alle som har støttet og hjulpet meg. Sånn faktisk.

Soli Deo Gloria

Contents

1	Introduction	1
1.1	Motivation	2
1.2	Goal of the thesis	4
1.3	Outline	5
1.4	Scope and delimitations	6
1.5	Contributions	6
2	Background	8
2.1	Taking inspiration from psychology	8
2.1.1	Awareness	8
2.1.2	The <i>Self</i>	9
2.1.3	Self-Aware Individuals	9
2.1.4	Public and Private Self-Awareness	10
2.1.5	Levels of Self-Awareness	11
2.1.6	Self-Aware Collectives	12
2.1.7	Self-Expression	12
2.2	Self-Awareness in already-existing Computing Systems	12
2.2.1	SA-related Computing Systems	12
2.2.2	Explicit self awareness in computing systems	16
2.3	Designing and Describing Self-Aware Computing Systems	17
2.3.1	Basis and Fundament	18
2.3.2	Reference architecture	19
2.3.3	Conceptual framework	20
2.3.4	Challenges & Limitations	21
2.4	Taking inspiration from natural phenomena	21
2.5	Oscillators and oscillatorsynchronization	22
2.5.1	Phase and frequency	22
2.5.2	Phase synchronization	23
3	Baseline	25
3.1	System target state: harmonic synchrony	26
3.1.1	Detecting harmonic synchrony	28
3.2	Phase synchronization	29
3.2.1	Problem statement	29
3.2.2	Synchronizing phases via phase adjustment	29
3.3	Phase and frequency synchronization	30
3.3.1	Problem statement	30
3.3.2	Synchronizing frequencies via frequency adjustment	31

4 Implementation and experimental setup	35
4.1 Simulator setup: the musical multi-robot collective	36
4.1.1 The simulator environment and its hyperparameters	36
4.1.2 The individual agent: a musical robot	37
4.1.3 Robot communication: the “fire”-signal	37
4.2 Synchronizing oscillator phases	42
4.2.1 Implementing and verifying Mirollo-Strogatz’s phase adjustment	43
4.2.2 Implementing and verifying Nymoen’s bi-directional phase adjustment	43
4.3 Synchronizing oscillator frequencies	44
4.4 Detecting harmonic synchrony	47
5 Experiments and results	50
5.1 Phase synchronization	50
5.1.1 Initial validation	51
5.1.2 Reproducing baseline results	51
5.1.3 Hyperparameter tuning	52
5.1.4 Comparing phase adjustment methods	54
5.1.5 Increasing degree of self awareness	55
5.2 Phase and frequency synchronization	60
5.2.1 Initial validation	61
5.2.2 Reproducing baseline results	61
5.2.3 Increasing degree of self awareness	64
6 Conclusions	68
6.1 General discussion	68
6.1.1 Conclusion	68
6.1.2 Ethical reflections	69
6.2 Further work	70

List of Tables

4.1	Notation and terminology used throughout the thesis for system components.	37
4.2	The environmental and collective hyperparameters present and used for the synchronization simulator in Unity.	38
4.3	The hyperparameters possible to alter for the individual musical robots in the synchronization simulator in Unity.	39
5.1	Hyperparameter setup for phase (ϕ) synchronization experiments.	51
5.2	Hyperparameter setup for ϕ synchronization baseline reproduction experiment.	51
5.3	Hyperparameter setup for ϕ synchronization hyperparameter tuning experiment.	52
5.4	Hyperparameter setup for comparative phase adjustment (Adj_ϕ) method ϕ synchronization experiment.	54
5.5	Hyperparameter setup for k_s nearest neighbours SA scope ϕ synchronization experiment.	57
5.6	Hyperparameter setup for collective size explosion ϕ synchronization experiment.	58
5.7	Hyperparameter setup for heterogenous SA scopes ϕ synchronization experiment.	60
5.8	Hyperparameter setup for phase and frequency (ϕ & ω) synchronization experiments.	61
5.9	Hyperparameter setup for first ϕ & ω synchronization baseline reproduction experiment.	61
5.10	Hyperparameter setup for ϕ & ω synchronization baseline reproduction and stabilization experiment.	62
5.11	Hyperparameter setup for second ϕ & ω synchronization baseline reproduction experiment.	63
5.12	Hyperparameter setup for k_s nearest neighbours SA scope ϕ & ω synchronization experiment.	65

List of Figures

1.1	Picture of fireflies flashing synchronously in a US National Park	1
1.2	Developed synchronization simulator inspired by synchronously flashing fireflies.	2
2.1	Previous and prevalent computational self adaptive agent frame- works.	13
2.2	Illustrations: two individual SoloJam example scenarios illus- trated differently.	14
2.3	Three Dr. Squiggles robots synchronizing to each other.	15
2.4	CoCoRo: the self aware underwater robotic AUV swarm.	17
2.5	Translated levels of self awareness from psychology to computer engineering and robotics.	18
2.6	Collective self awareness schema.	19
2.7	A rererence architecture for a self aware, and self expressive, com- puting system.	20
2.8	Sketch of Miroollo & Strogatz's "standard" phase adjustment.	23
3.2	Frequency spectrograms illustrating the absence and presence of harmonics and overtones in audible waveforms.	27
3.3	Nymoen et al.'s harmonic synchrony detection schema.	28
4.1	Schema of the developed musical multi robot collective.	35
4.2	Editing a longer and pure harp pluck to become a shorter fire signal.	41
4.3	Schema of frequency based fire signal assignment.	42
4.4	Illustration of Miroollo & Strogatz's "standard" phase adjustment (Adj_ϕ) method.	43
4.5	Illustration of Nymoen's bi-directional phase adjustment (Adj_ϕ) method.	44
4.6	Simulation run plot: frequency plot.	45
4.7	Simulation run plot: harmonic synchrony evolution plot.	49
5.1	Experiment results for ϕ synchronization baseline reproduction experiment.	51
5.2	Experiment results for ϕ synchronization hyperparameter tuning experiment.	53
5.3	Experiment results for comparative phase adjustment (Adj_ϕ) method ϕ synchronization experiment.	55

5.4	Experiment results for ϕ synchronization k_s SA scope tuning experiment.	57
5.5	Experiment results for ϕ synchronization d_s SA scope tuning experiment.	58
5.6	Experiment results for first ϕ & ω synchronization baseline reproduction experiment.	62
5.7	Experiment results for ϕ & ω synchronization baseline reproduction and stabilization experiment.	63
5.8	Experiment results for second ϕ & ω synchronization baseline reproduction experiment.	64
5.9	Experiment results for ϕ & ω synchronization k_s SA scope tuning experiment.	65
5.10	Experiment results for ϕ & ω synchronization d_s SA scope tuning experiment.	66

Chapter 1

Introduction

Taking inspiration from the fascinating natural phenomena of self synchronizing fireflies, as e.g. seen in Figure 1.1, a synchronization simulator imitating and modelling this process for—instead of fireflies—a collective of musical robots, as e.g. seen in Figure 1.2, has been designed and tested.

The effects, particularly in terms of the ability to and performance of synchronization, is measured and evaluated when experimentally altering the properties of the musical robot collective as a whole (as e.g. collective size), or individual musical robots's hyperparameters (e.g. the number of neighbouring robots each robot listens to for self synchronization, later referred to as the robot's *self awareness scope*).



Figure 1.1: Synchronous fireflies at Congaree National Park, United States.
Photo¹: @columbiasc & @_flashnick.

¹<https://avltoday.6amcity.com/synchronous-fireflies-congaree-national-park/>
(accessed 2022.05.17)

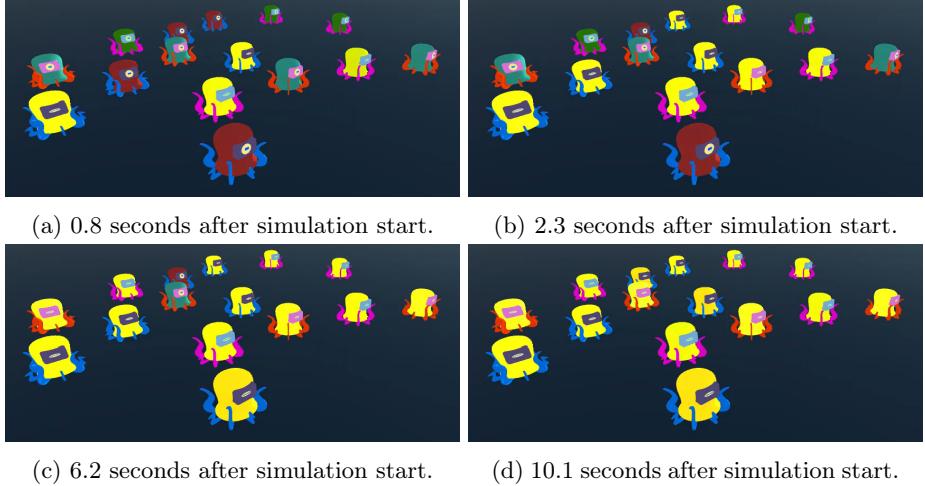


Figure 1.2: A typical simulation run of a musical multi robot collective (here of size 15) entraining to achieve harmonic synchronization in the developed system. As we can see², the robots start flashing yellow and blinking with their eyes asynchronously at first in 1.2a, but end up synchronized in 1.2d. Due to our special synchronization type being *harmonic synchronization* (cf. 3.1), as well as an implementation choice of robots only “flashing” or “firing” every other phase climax (i.e. $\phi = 1$), musical robot collectives do not in fact have to all fire simultaneously in order to be harmonically synchronized, but all robots firing simultaneously is also a validly harmonically synchronized robot collective.

1.1 Motivation

Engineering a computing system for a certain environment often requires some knowledge of said environment; both on the end of the system-creator, as well as for the computing system in turn. This knowledge will then have to be used in the environment the computing system is situated in, in all the situations and possible states the system can be in during its lifetime. However, foreseeing and predicting, at design-time, all possible future states and scenarios a computing system will be in during its lifetime is often hard, and sometimes impossible (in the case of coincidental faults e.g.). Increasingly extreme, complex, dynamic, and ever-changing environments, simply enlargens this challenge.

Extreme challenges and physical barriers within communication between decentralized and mobile robots and its human operator, like too large latencies, bottlenecks presented by using the same limited bandwidths, short possible ranges, or the inability to use global satellite systems e.g. in underwater vehicles [23], leads to the necessity of enabling systems to autonomously and online control themselves and perform missions without communication from remote operators like humans. This calls for online and continuous learning at run time, where computing systems are themselves able to observe, learn, adapt, and act on their own, independently from their creator; *autonomic computing*

²See video at <https://www.uio.no/ritmo/english/projects/modeling-and-robots/media/phasesyncrecord.mp4> (accessed 2022.05.27)

that is. Hence, technological development and research has e.g. within under-water robotics gone from enabling remotely operated vehicles (ROVs) in the 1980s (for e.g. oil and gas exploitation at depths unreachable to human divers), over to more self adaptive and autonomous underwater vehicles (AUVs) closer to in the 21st century—the demand of which is expected to grow 37% in the 2018-2022 period [23]. As a result, the underwater robots can in this instance thus, due to the increased degree of autonomy and consequently implicit reduction in communication need, travel on its own on even longer and extensive missions without the need for a live connection to any human operator—e.g. on seabed mapping missions, underwater machinery or structure maintenance, as well as seabed cleaning.

However, and especially within collectives of multiple robots [17, 27], communication between individuals—even though they might be autonomous individual robots—is still a challenge one has to deal with when designing them. Coordination, for humans and robots alike, often tend to demand a lot of communication. To add on top of this, when networks or collectives of sub-systems like wireless sensors are communicating with each other, they often have their own internal clock which often do not align. Hence, we see that both the magnitude of communication needed, as well as the inherent challenges with communication, call for effective approaches to communication with coordination as the goal.

As K. Konishi & H. Kokame [11] points out, one technical problem in e.g. wireless sensor networks is that each sensor node rely on accurate internal and synchronized clocks, especially from the perspective of sensor fusion and co-ordinating communication among nodes. Thus, time synchronization becomes recognized as one of the crucial problems in distributed and wireless systems [28]. Various attempts at synchronizing messages and communication, in order to e.g. align messages exchanged at unsynchronized intervals, have been made [28]; however, such attempts and protocols often require the computation of message exchange and processing, which wastes the limited computation capability of nodes and causes communication delays. To this, Konishi & Kokame [11] also point out how for synchronized pulse coupled oscillators (PCOs), such computation is not required. If internal clocks in communicating nodes are instead altered and synchronized to other nodes's clocks—instead of letting all clocks stay unsynchronized and constantly trying to compensate for this—it is rather apparent that getting rid of this need for compensation will considerably reduce computation needs while synchronized communication is still achieved.

Thus, by achieving time synchronization [28] through synchronizing pulse coupled oscillators, both the computational load, in addition to the need for communication, is considerably reduced; hence freeing up and saving valuable resources like e.g. processing power and energy (battery) in networks of individuals and possibly autonomous nodes, agents, or robots.

On another hand, self-awareness concepts from psychology have given inspiring new approaches for engineering computing systems which operate in complex dynamic environments [14]. As computer scientists and roboticists have recently brought more focus on the field of self awareness relatively recently [13, 14], many new conceptual tools and frameworks have been developed. As such, it is not necessarily more intelligence or self awareness capabilities having been brought into new systems; no, on the contrary, for also older and already existing computing system one now has the ability to describe and discuss said

computing system's self awareness capabilities—even if it was not originally designed with any such capabilities in the first place. E.g. K. Glette et al.'s translated levels of self awareness, from the man called “the father of cognitive psychology” (Ulrich Neisser), into computing systems [14], or P. Lewis et al.'s reference framework for computationally self aware systems [4], give rich descriptive power for both designing novel computational systems, as well as describing already existing ones—as one inherently within the system include the aspects and notions of self awareness, instead of e.g. simply adding a layer of attention on top of the system as has been custom to do before.

With self awareness, online and continuous learning is achieved to a higher degree in contrast with other approaches (like ODA and MAPE-K), due to the limitations and downsides of these older approaches, as well as the advantages and upsides to considering computational self-awareness in computing systems. If one wants to achieve continuous adaptation of a system or of system-components (e.g. in a collective) — some sort of intelligence might be necessary to endow it with. Endowing computing systems with Self-Awareness can be beneficial in several respects, including but not limited to a greater capacity to adapt, to build potential for future adaptation in unknown environments, and to explain their behaviour to humans and other systems [13].

Explainability has with time only become more and more relevant, also within artificial intelligence (AI) (hence the popularity of the term *explainable AI*), as increasingly autonomous and automatically systems are making real life decisions with serious consequences increasingly on a day to day basis. L. A. Dennis and M. Fisher present explaining and verifiable agent-based systems, where rational agents who possess goals, beliefs, desires, intentions e.g. make decisions that in their opinion should be able to be questioned and giving an answer when prompted [5]. Self awareness enables such questioning and answering by autonomous agents, as the agents themselves, since they are themselves aware of their knowledge, thoughts, goals, desires etc., can explain what lead to their actions. Or in other words, questions like “what type of self awareness [14] (knowledge) lead to a certain self expression [14] (action)” can be answered by self aware and self expressive agents. Such an example can be read by P. Lewis et al. [13] (3.3.4) when they summarize their novel computational self awareness framework by enabling self aware computing systems to produce sentences like the following: Peter (span) is aware of Ada's (scope) goal (aspect) to reduce the power usage (object). Further, Ada (span) is aware of her own reasoning (meta-self-awareness) about what to do (act) about it.

Endowing computing systems with Self-Awareness can be beneficial in several respects, including a greater capacity to adapt, to build potential for future adaptation in unknown environments, and to explain their behaviour to humans and other systems. As we can see in various Music Technology Systems, this endowing can also give rise to interesting cooperative and coordinating behaviour.

1.2 Goal of the thesis

In this MSc-thesis, we will explore an exciting and relatively new translation of the concepts and notions regarding self-awareness — as they pertain to humans and animals especially — from the domain of Psychology, into the domain of computation and engineering. The problem to explore will mainly consist of

studying the effects differing self-awareness levels, varying collective-sizes, levels of task difficulty (like more complex behaviours, and limited communication) - have on usefulness, system dynamics, overall performance, more intelligent systems, and scalability (in this case specifically within a musical robot system) when performing a task. Whether the effect of endowing the musical robot system with increased computational self awareness will be better decision making (or *self expression* [14]), versatility and ability to adapt quickly and online in a rapidly dynamic and everchanging environment, as they continually learn by themselves further handling continuous environmental change, or not will be investigated.

Generally, the aim is to discover the effects of endowing computational systems/robots with self-awareness capabilities / abilities. More specifically, the aim of the thesis is to explore and investigate whether—and to which degrees—increased degree or level of computational self awareness in musical robots lead to increased performance in a collective musical task to be achieved; namely, achieving synchronization (*harmonic* at that, cf. 3.1)—much like and inspired by fireflies in nature adjusting and synchronizing themselves and their flashing to each other.

With that, the following research questions is thus introduced, with the hope of answering them in the thesis:

Research Question 1:

To what extent will increased levels or degrees of computational self awareness in individual musical robots lead to the musical collective at large being able to achieve its collective musical task—being (*harmonic*) synchronization—faster compared to with lower levels or degrees of computational self-awareness?

Research question 2:

To what extent can harmonic synchronization performance in musical robot collectives be maintained despite lower degrees or levels of computational self awareness? Are in fact global connections in oscillators necessary to maintain harmonic synchronization performance, or can more relaxed and less constrained connections in oscillator networks still give rise to similar synchronization performance?

Research Question 3:

In what ways do computational systems—specifically a musical robot collective—exhibit and display self awareness (and corresponding self expressive) capabilities, compared to in humans and animals?

1.3 Outline

The master's thesis is structured as follows. Firstly, past and related work relevant for the thesis will be presented in Chapter 2, laying the theoretical and inspirational foundation this thesis is basing itself upon.

Then, in Chapter 3, a particularly important and foundational approach to oscillator synchronization is presented in detail, serving as the starting point and baseline for the synchronization method the musical robots in our developed

system will utilize in order to achieve the target state also presented and achieved by the authors of the baseline approach.

After the main specific approach of comparison and inspiration is presented in Chapter 3, the specifics and details of the newly designed and implemented synchronization simulator will be expounded in Chapter 4. Here, experimental setups will also be shown, for which we want to evaluate synchronization performance for in Chapter 5. In this chapter, we will also analyze the experiments results from the motivated and set up experiments.

Finally, results found in Chapter 5 will be discussed and put into a broader context, as the thesis is concluded and discussed on a higher level, summarizing the whole thesis. Before references and eventual appendices are shown, possible future work as well as the main reasons why they were not possible to explore in the thesis work will be given.

1.4 Scope and delimitations

No training of any neural networks or AI models was performed to achieve synchrony in the newly developed synchronization simulator; so far no machine learning is used.

Even though the evolution of how musical robots progress and evolve towards a state of synchronization, no evolutionary algorithms or evolutionary searching algorithms are implemented or tested in this thesis either.

Instead, as is ubiquitous in and very typical for multi agent systems and swarm robotics applications, the phenomenon of fairly simple rules endowed in rather simple agents / robots leading to an emergent and collective effect is observed—also in the newly developed synchronization simulator in this thesis. For this thesis's sake, this emergent and collective effect is musical robot collectives reaching a state of *harmonic synchrony* (cf. 3.1).

1.5 Contributions

The main contribution this thesis work has led to is the design, implementation, and testing of a novel synchronization simulator for musical robots modelled as oscillators, further explained in 4.1, achieving time synchronization—previously having been recognized by scientists and roboticists to be one of the crucial problems in distributed and wireless systems.

In said synchronization simulator, synchronization methods (for both oscillator phases ϕ and frequencies ω), collective properties of the musical robots, as well as individual properties and hyperparameters in musical robots (hence enabling heterogenous robots), can all be experimentally altered. The resulting synchronization process, including all the musical robots's interactions and entrainment towards harmonic synchrony, can then be visually seen and heard in real-time, as well as thoroughly analyzed using simple plotting scripts after simulation runs have ended. Hence, a rich exploration space—more than able to be further expanded and built upon—and testbed for achieving (*harmonic*) synchronization is opened up to all that are interested in exploring it and have a GitHub account. Furthermore, the synchronization results and performance of said explorations are easily visualized and heard. This then democratically

enables for the open access to experimentation with creative and novel synchronization methods in oscillators, in order to seamlessly qualitatively and quantitatively assess their efficacy and efficiency at achieving harmonic synchrony.

Secondly, and differently to earlier approaches synchronizing pulsed coupled oscillators (harmonically at that [22]), the synchronization system developed in this thesis in Unity does not limit oscillator frequencies to stay within a certain range, and oscillator frequencies can both be adjusted to be larger than the max initial frequency, as well as smaller than the minimum initial frequency. In e.g. Nymoen et al.'s firefly inspired oscillator system, frequencies ω are limited to staying within the range of [0.5Hz, 8Hz].

Chapter 2

Background

In this chapter, pointers to and clusters of recent work in the fields relevant to the thesis, especially where they intersect, are presented. These give a sense of where the thesis is coming from, and what lead to the explorations made later on.

2.1 Taking inspiration from psychology

Self awareness (SA) in psychology:

The notion and idea of *self awareness* has been of great interest to humans throughout history. And despite its sometimes abstract and elusive nature, many thinkers (like e.g. psychologists and philosophers) have previously tried to give an explanation and description of this phenomenon or capability, primarily as observed in individual humans and animals. One has also tried to extend these descriptions and developed concepts to explain collectives consisting of individuals. A good survey of such descriptions and explanations regarding Self-Awareness in Psychology is given in some recent books on Self-Aware Computing Systems [13, 14] by P. R. Lewis et al, and will be used as a basis for this section and essay.

The notion of self awareness might give associations to other psychological terms like perception and consciousness, and we will look at some relationships these have with the main notion here of self awareness, in the psychological and (briefly) philosophical literature.

What exactly is self awareness? Is it *awareness* of a *self*? If so, this leads to the question: what constitutes a *self*? And also, what does really *awareness* mean? Some answers given in psychological literature, and in later time taken inspiration from into computer science, will now be explored.

2.1.1 Awareness

Before Lewis et al. [14] inclusively and succinctly introduces the concepts and interpretations relating to Self-Awareness, as they pertain to the psychological

literature especially, they start off by giving the definition of *awareness* by The *Oxford English Dictionary*, which is "knowledge or perception of a situation or fact." Hence (by simply adding the prefix *self*), "situational or factual knowledge/perception about a *self*" might seem to be a reasonable first intuition of the notion of Self-Awareness. We will now discuss further what might constitute a so-called *self*.

2.1.2 The *Self*

Prominent western philosophers, like David Hume and Immanuel Kant, had some differing views on what constitutes a *self* — although both might seem reasonable and relevant.

As Lewis et al. describes it [14], Hume is known for postulating that all human knowledge is induced from experience. This then means for Hume that the *self* is no concrete and physical entity in the world, but rather the bundle of experiences or perceptions unique to an individual.

On the other hand, for Kant, the scope of the self is significantly extended — as he argues for the self being an actual entity being the subject of experiences, which is common through space and time. This subject or *self* then combines information from experiences with concepts in the mind, as well as with the imagination — as well as acting upon this combined knowledge.

2.1.3 Self-Aware Individuals

Given a Humean view of the self (as described above), one might then consider self awareness to consist of an individual's knowledge of its experiences. Whereas the Kantian view of the self yields a self awareness consisting of a distinct and physical individual subject to experiences, combining information from these experiences, concepts, and the imagination.

Early appearing terms in the psychology literature, pertaining to Self-Awareness and notions of self — being reminiscent of these two differing views by Kant and Hume — are found in the distinction between the *objective/explicit self*, and the *subjective/implicit self*, firstly elucidated by W. James in 1890 [29], and further expounded in detail by Duval [7]. Here, the *implicit* self, often referred to as the self-as-subject, is the "me"-self, subject to experiences. The self awareness in this "me"-self then, according to Duval, is described as "*a state of consciousness in which attention is focused on events external to the individual's consciousness, personal history, or body*" [7]. The *explicit* self however, often said to be the "I"-self, describes a self-as-object which can be discerned, observed, and reasoned about in the world. This "I"-self then, according to Duval, has an objective self awareness described as being "*focused exclusively upon the self and consequently the individual attends to his conscious state, his personal history, his body, or any other personal aspects of himself*" [7].

One well-known "self awareness-test", having been used to deem whether someone or something is self-aware or not, is the so-called *mirror test*. This consists of observing an individual's capability (or lack thereof) of noticing or self-recognizing, in a mirror, a discreet change in their own appearance (that was caused by someone else typically, while they were not paying attention e.g.). The presence of such a capability is argued (e.g. by Asendorpf et al. [1]) to

demonstrate the presence of a *secondary representation* (a constructed conceptual model if you will) of oneself and the world (the primary representation being the real world).

Given that the ability to construct such a conceptual model of oneself is required in explicit self awareness, one could wonder whether passing this "mirror test" would then be satisfactory for someone or something (like a system) to be self-aware. Haikonen [9] however demonstrated how easy it was for a not-too-sophisticated computing system to pass this test — and hence weakened the validity of this "mirror test" to definitively conclude whether someone or something was indeed self-aware or not.

As should be more and more apparent, there has been — and still is — a lot of ongoing discussions about what should be considered self awareness, and what should not — as well as about all the different variants observed (by e.g. Morin [18]).

Morin defined self awareness as different from, but building upon the previously mentioned consciousness (where this in some cases are deemed more "primitive" aspects of self awareness), in that "self awareness is the capacity to become the object of one's own attention." However, he doesn't deny the subjective aspect of the self-aware individual, in that he believes a self-aware organism to be one that "becomes aware that it is awake and actually experiencing specific mental events."

The importance of including such a *subjective* (perhaps more underlying and primitive) level of self awareness, lies in the realization that self-aware agents (being uniquely situated in the world, collecting unique data from their unique sensing-apparatus) **is** subjective in nature. As such, since possibly heterogenous agents might collect widely different data (due to their situation and sensing capabilities) from the same exact phenomenon, it is crucial that we know from which *subject* the data we have originates from.

Perhaps or perhaps not due to this, others would object to Morin's definition, and actually include this subjective and implicit (or in other words "perceptual (or pre-conceptual)", as Lewis et al. [14] put it) experience in their definition of self awareness — leading to self awareness and consciousness being overlapping concepts. This would also lead more complex and "full-stack" Self-Aware systems to being extensions of this simpler and more primitive level of Self-Awareness.

2.1.4 Public and Private Self-Awareness

The distinction, as discussed above, between the implicit/subjective and the explicit/objective self and corresponding self awareness has given rise to many further developments in this notion, as is presented by Lewis et al. [14]. However, these further developments separate this seeming distinction of *private self awareness* (previously described as e.g. personal, historical, and bodily knowledge regarding the self-as-object) and *public self awareness* (previously described as awake, subjective, and unique experiences external to the self) a bit differently.

Rather than the above notions of public and private self awareness, **private self awareness** instead here refers to the obtaining of "knowledge of internal phenomena, typically externally unobservable and accessible only to the individual" [14]. On the other hand, **public self awareness** now concerns the

gathering of knowledge based on phenomena external to the self. If we choose to look at public self awareness in the implicit/subjective form, we have a "minimally" self-aware individual who is merely able to perceive (subjectively) the external environment in which it is situated. Conversely, we could also have a publicly self-aware individual in the explicit/objective form — which would be more concerned with public appearance, and how the individual is perceived in its environment (e.g. by other individuals, in social settings). These definitions of private and public self awareness will be used from now on.

2.1.5 Levels of Self-Awareness

There is large consensus about, due to the many degrees or variations to which an individual can be self-aware, that self awareness is not *binary* or singular — but rather on a spectrum. As Lewis et al. highlights ([13, 14]), varieties of an individual's self awareness capabilities are often associated or characterized with one or more levels of self awareness. As such, several have attempted at defining such self awareness levels.

One such definition, being comparatively broad and wide (to other similar treatments), containing 5 increasingly complex self awareness levels, is given by Neisser [21]:

1. Ecological self

The ecological self is the most minimal form of self awareness. It permits sufficient knowledge only for basic stimulus-response behaviour, as the individual has a basic awareness of stimuli. The ecological self can be thought of as the minimum requirement for the individual to not be unconscious.

2. Interpersonal self

The interpersonal self enables the individual to possess a simple awareness of its external interactions, permitting limited adaptation to others in the performance of tasks.

3. Extended self

The extended self extends the interpersonal self to permit reflection of interactions over time. The individual is aware of the existence of past and future interactions.

4. Private self

The private self allows the individual to process more advanced information concerning itself, such as thoughts, feelings and intentions.

5. Conceptual self

The conceptual self, or self-concept, is the most advanced form of self awareness, representing that the individual is capable of constructing and reasoning about an abstract representation of itself.

This final level, the **Conceptual self**, is very much connected to the much-discussed term of *meta-self awareness*. Meta-self awareness refers to the awareness of someone that they themselves are self-aware. This includes reasoning

about its own self awareness, and is argued to be necessary (at least in humans) in order to adapt to changes, and to direct ones attention where it is needed — and that without it, thoughtlessness would ensue.

We will see later on in Section 2.3 how these levels come into play in a fundamental basis and a proposed reference framework, for describing Computational Self-Awareness.

2.1.6 Self-Aware Collectives

In her comparison of three distributed and collective biological systems (the brain, the immune system, and ant colonies), Mitchell [16] elucidates how — not only individuals, as has been the case until now, but — also collectives (consisting of individuals) can exhibit apparent global self awareness properties/capabilities. As she goes on to explain and make her case, she further elucidates how these apparent self awareness properties are only seen globally — but not in individuals (which in this case are neurons, ants, or lymphocytes - also known as white blood cells). The explanation for this, she proposes, is due to the distributed and statistical nature of this apparent self awareness. Through the many individual interactions in the collective — self awareness gets built up *bottom-up*, and then gets fed back down to the lower-level components in a *top-down*-fashion.

Given that this apparent self awareness then was not present in the collective's individual components, but in the global collective, as a result of the many interactions of the lower-level components — we might say that the collectives exhibit **emergent** self awareness.

Hence, we can both analyse and describe Self-Awareness capabilities in individuals, as well as collectives of individuals.

2.1.7 Self-Expression

Lewis et al. [14] argues that self awareness properties alone (i.e. the ability to continuously acquire, update, and synthesise the self's models and knowledge about itself and its experiences) is of limited value, unless it is accompanied by associated behaviour. As Chandra et al. [4] argue, later on in the same book — if the self only takes data in, but does not use that data to do anything, then it is essentially a "data-sink". Hence, the accompanying notion of *self-expression* is introduced as "behaviour based on self awareness" [14]. This behaviour could be highly dynamic (e.g. enabling self-adaptation, self-reconfiguration, self-explanation), as well as constituting a more "standard" system behaviour, where self awareness is not considered at all.

2.2 Self-Awareness in already-existing Computing Systems

2.2.1 SA-related Computing Systems

2.2.1.1 Self adaptive systems

Run-time adaptation in Computing and Engineering systems is not a new thing, and this has been of interest across a range of research communities for quite

some time [4]. The reason for this becomes apparent when one considers the problem of tackle unforeseen (at least at design-time) operational changes that can occur during run-time (like e.g. faults).

An agent — in the context of computing systems — can be said to be a computing system situated in an environment (which it can perceive), having some goals, and possible actions it can perform (through actuators) to reach its goals. It can also be informed by some knowledge, and/or directly the signals it gets from its sensors.

Especially within the engineering of **Self-adaptive** Systems and Intelligent Agents, such agents are often designed with properties which can be called Self-Awareness and -Expression properties — even though these specific terms are not used.

Notable examples in accordance with the autonomic computing paradigm ([4]) comprise for example ODA (Observe-Decide-Act) loops, originally stemming from the OODA (Observe-Orient-Decide-Act) loop in the military domain. As the acronym suggests, systems implementing the ODA-loop continually Observe (their environment and potentially themselves through sensors), Decide (planned future actions based on current or future states), as well as Act (on one of those previously planned actions). A schematic of such an ODA-loop can be seen in Figure 2.1.

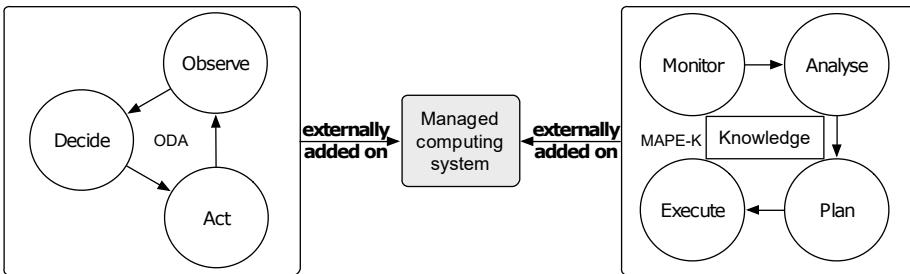


Figure 2.1: Figure of the Self-adaptive Systems ODA (left) and MAPE-K (right). We also see the externality of the attempts at implementing meta-self awareness in these systems, in that they simply *add on* an external Managed computing system (which is monitoring the rest of the system). Figure reused from Chandra et al. [4].

Another notable Self-adaptive System we can see in Figure 2.1, is the MAPE-K architecture (an acronym for Monitor, Analyse, Plan, Execute, and Knowledge). The great and groundbreaking benefits of implementing e.g. ODA-loops and the MAPE-K architecture, performing online observation and knowledge-gathering (hence having endowed their systems with some sort of self-knowledge/self-awareness) allowing for online self-adaptation — have been greatly reaped. However, these architectures have not explicitly considered the psychological origins of self awareness, and lack some benefits compared to the introduced frameworks and architecture in Section 2.3.

2.2.1.2 Music technology systems and musical robots

As particularly relevant to the thesis project, endowing Music Technology Systems with Self-Awareness is of great interest. There have been several such

approaches, particularly within Interactive and Active Music Systems. One example includes:

SoloJam: Achieving de-centralized participation (by co-ordinating and co-operating agents) in a coherent and interactive music system

Historically, the creation and production together with the listening and perception of music has been (and still is) relatively "one-way'ed"; in that the musicians/artists produce the music, and the fans/listeners consumes the music. A growing interest in interactive music (like Guitar Hero etc.) has been seen the last decades. The authors of the music system SoloJam [3] want participants with little or no musical experience and training to be able to play independently and decentralized (from say an expert) music as a whole, by each playing solos made from their own devices. Here each of their solos will be slight variations of the previous participant's solo. See illustration in Figure 2.2.

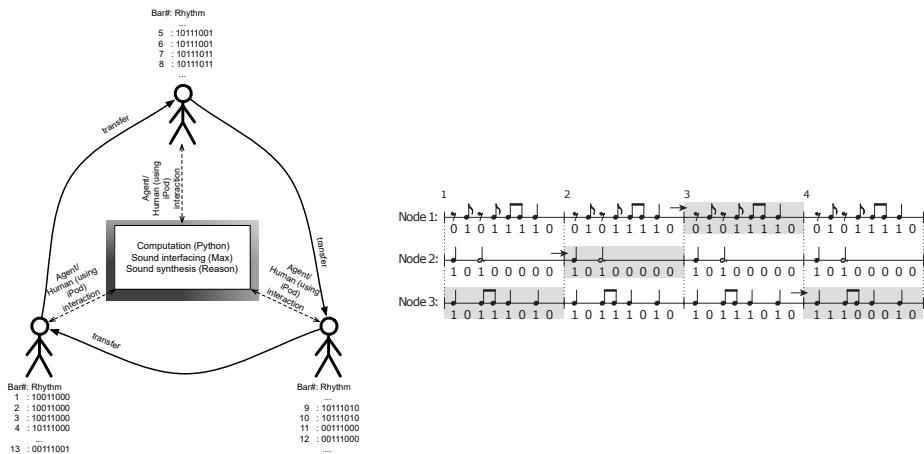


Figure 2.2: Illustrations: two individual SoloJam example scenarios illustrated differently.

The solution to how this decentralized, coherent and interactive music system consist of elements inspired by the Economical Sciences. Namely:

1. Auctions (specifically the Vickrey auction)
2. Utility (profitability/deservedness measure)

The function for utility is given by the expression:

$$u_i = \frac{c}{(1 + aD_l)(1 + bT_l)} \quad (2.1)$$

It is essential with a properly designed utility function for the system to behave coherently e.g. This utility-function is inversely proportional with *a*) differences between bidder and auctioneer's/leader's rhythm patterns D_l , and *b*) the time/#bars the leader has played its rhythm pattern T_l - as well as constrained

by and applied to a couple of other clauses (to ensure some wanted/avoid some unwanted effects).

"Shaking"-/"moving"-patterns (as registered/sensed by the devices' inertial units) are translated into rhythm patterns - which then are calculated a **utility score** (by the function as mentioned above and outlined in more detail in the paper) for, used as bid in an auction.

The winner of the auction will (when considering knowledge of the leader-rhythm-pattern, as we do when $a = 1.0$ instead of 0.0) be those rhythm patterns which are the closest rhythm-patterns (i.e. not completely equal) to the leader's rhythm-pattern (in terms of the Hamming-distance) - assuming e.g. that the participant didn't just play before another participant, nor that it had played its rhythm-pattern for a very long time. In the paper, we then notice how the authors study how **varying degrees of self awareness** ($a=1.0$ versus $a=0.0$) affect the overall system dynamics — then in the form of musical coherence, as well as decentralized circulation.

Musical robots: the Dr. Squiggles

The Dr. Squiggles robots, which are M. J. Krzyzaniak and RITMO's musical robots, as e.g. can be seen in Figure 2.3 as they synchronize to each other, have been used in several music technology systems and projects¹ previously.



Figure 2.3: Three Dr. Squiggles robots synchronizing to each other. Photo: from M. J. Krzyzaniak's video *Squiggles Equilibrium*²

SoloJam Island

Combining the two musical systems presented above, Pierre Potel implemented islands of Dr. Squiggles robots in a SoloJam scenario, and hence naming the novel musical robot system *SoloJam-Island*³. Corresponding 3D-models of the Dr. Squiggle robots, for said simulation, were designed by Pierre Potel.

¹<https://www.uio.no/ritmo/english/news-and-events/news/2021/drsquiggles.html>
(accessed 2022.05.17)

²<https://www.youtube.com/watch?v=yN711HXPfuY> (accessed 2022.05.27)

³<https://github.com/67K-You/SoloJam-Island> (accessed 2022.05.17)

2.2.2 Explicit self awareness in computing systems

The extensive interest of studying self awareness in natural fields like psychology, sociology, and philosophy might not come as a surprise, as self awareness primarily is found and exhibited in humans and animals. However, there has recently been an uptake in interest for explicitly studying self awareness in computing systems.

Recently, (often collaborative) efforts within the computer science & engineering community [12, 13, 14] have lead to a greater systematic understanding of self awareness, and how it can be used and evaluated within Computing Systems.

Some of these fields, now explicitly interacting with the notion of self awareness — as they were presented in overviews given by P. R. Lewis et al. [13] and S. Kounev et al. [12] — include autonomic computing, machine learning and artificial intelligence, multi-agent systems, self-organizing and self-adaptive systems, situation-and context-aware systems, reflective computing, model-predictive control, as well as work from the models@run-time community. But also, fields like Robotics, complex IT systems, computer engineering, and reflective architecture [13].

Some quite notable research projects and initiatives within computer science and engineering, explicitly engaged with the notion of self awareness in computing has been performed (like the SEEC project at MIT and University of Chicago, the ASCENS and EAssets/pics FP7 EU Projects, as well as the SEAMS Dagstuhl Seminars and workshop series) [12].

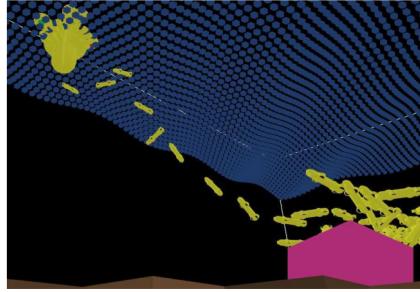
2.2.2.1 Specific Examples

Cognitive Radios

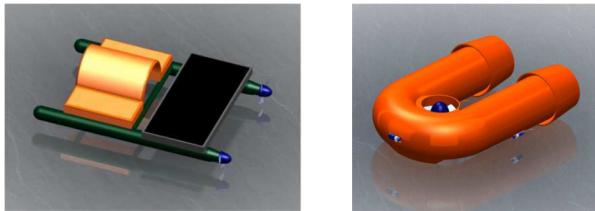
A — perhaps simpler or minimal — specific instance of a computing system where Self-awareness have been interpreted and applied, is the so-called *cognitive radio*. The devices in this system control their own capabilities and monitor other devices' capabilities through communication with these other devices. This "*enables them to improve the efficiency of communication by negotiating changes in parameter settings*" [13]

CoCoRo: the self aware underwater AUV swarm

More complex self-aware computing systems have also been explored. As demonstrated by Mitchell [16], also T. Schmickl et al. considers self awareness in a collective system — an underwater swarm robotic system [27] in this case (see Figure 2.4). This robot collective shall in this scenario search the sea-bottom for specific objects of interest (like toxic waste, but also sunken objects e.g.). Here, the collective/swarm is supposed to e.g. be self-aware of its own state (of whether it still is searching, or if it has found a target at the bottom of the sea) — in order to either stay together searching, or communicate to a second (surface level) swarm that an object of interest is found. Like in Mitchell's case, these autonomous underwater vehicles (AUVs) does have emergent self awareness properties that the individuals themselves do not possess. Therefore, as the authors say in [27], "*Finding such targets will only be efficient if the AUVs work together.*"



(a) Swarm searching the seabed



(b) Basestation

(c) AUV

Figure 2.4: CoCoRo: the self aware underwater robotic swarm consisting of AUVs alternating between completing missions at the seabed and returning to the surface level basestation. Figure reused from T. Schmickl et al.'s paper [27].

Self-Aware and Self-Expressive Camera Networks

Lukas Esterle et al. [26] has demonstrated how the fundamental concepts and building blocks of Computational Self-Awareness and Self-Expression can handle resource-critical- as well as communication- and utility-tradeoffs, maintaining network topologies, performing distributed object detection, performing tracking handover, as well as achieving more Pareto-efficient outcomes. They also believe these capabilities are not exclusive to camera networks — but are in fact confident these technologies (of computational SA and SE) can enable "*other types of systems and networks to meet a multitude of requirements with respect to functionality, flexibility, performance, resource usage, costs, reliability, and safety*" [26].

Increased local neighbourhood self awareness J. Cao et al. [2] found that as they increased agents's degree of *local neighbour awareness* in a multi-agent reinforcement learning (MARL) scenario, they obtained higher collective performance and cooperation playing a challenging video game (StarCraft II).

2.3 Designing and Describing Self-Aware Computing Systems

Psychology has served as the basis & main source of inspiration. Here we will look at a translation of Self-Awareness, from Psychology to Computing. We will hence look at an introduction of the new notion of *Computational Self-Awareness*.

Translating concepts from the Psychological field into the Computational field gives rise to rich sources of inspirations (for engineers and researchers who might not have thought psychologically about designing and describing computing systems), as well as hopefully clear, structured, principled conceptual tools to in a consistent manner describe and design so-called *Self-Aware Computing Systems*.

2.3.1 Basis and Fundament

Lewis et al. [14] have introduced a basis and fundament for Computational Self-awareness and Self-expression, that is founded in the psychological literature on Self-Awareness, as presented in Section 2.1.

More specifically, they base their new notion of computational self awareness on computationally translated versions of the key self awareness concepts discussed in Subsections 2.1.4, 2.1.5, and 2.1.6, as well as 2.1.7.

First off, public and private self awareness now considers computational systems, instead of human selves — but these definitions are still the same (only that now the *self* is the *computational system*).

Secondly, Neisser's five levels of self awareness is now translated by P. Lewis et al. [14] into five corresponding levels, as shown in Figure 2.5 below.

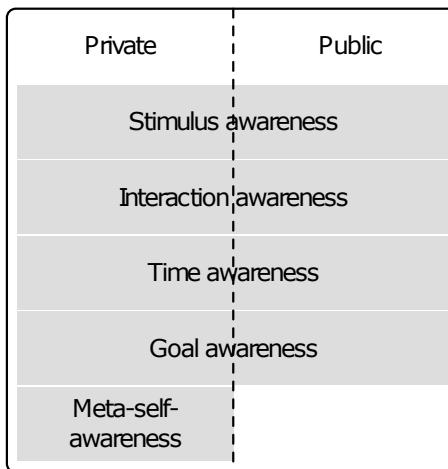


Figure 2.5: Translated levels of self awareness from psychology to computer engineering and robotics. Figure credits: Lewis et al. [14].

Thirdly, computational self awareness can both be considered in individual agents/computing systems, or in arbitrary agent-collectives (consisting of several sub-collectives) — as seen in 2.6 below.

Lastly, *Computational Self-Expression* is similarly defined as (only now for a computing system as the *self*) **behaviour based on computational self awareness**.

Kounev & Lewis et al. also gives a definition of Self-aware Computing [12]:
Definition 1.1 Self-aware computing systems are computing systems that:

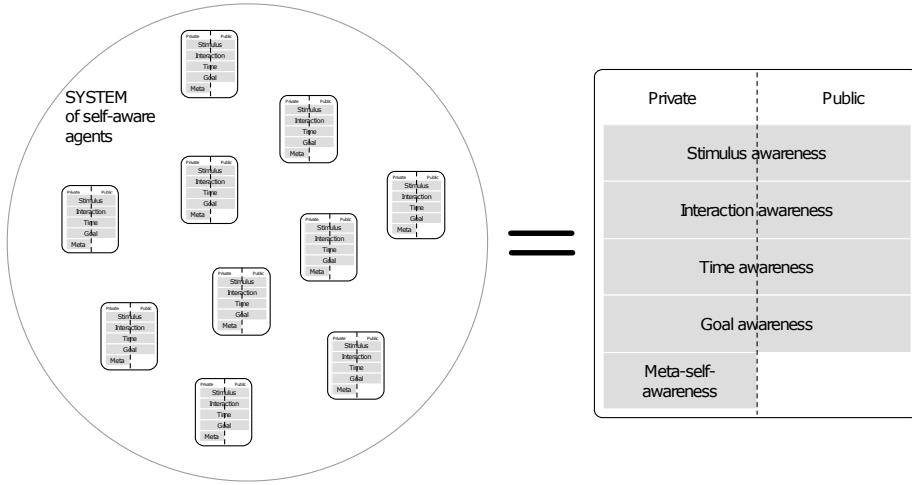


Figure 2.6: Collective self awareness schema, as entities can be seen as consisting of many sub-entities. Figure credits: Lewis et al. [14].

1. *learn models* capturing *knowledge* about themselves and their environment (such as their structure, design, state, possible actions, and runtime behavior) on an ongoing basis and
 2. *reason* using the models (e.g., predict, analyze, consider, and plan) enabling them to *act* based on their knowledge and reasoning (e.g., explore, explain, report, suggest, self-adapt, or impact their environment)
- in accordance with *higher-level goals*, which may also be subject to change.

The last part of point 1. in Kounev & Lewis et al.'s above definition of self aware computing systems, mentioning "ongoing basis", is key. It is not enough for a computing system to simply possess data or pre-programmed knowledge to be considered self aware; the computing system needs to be able to collect, analyze, and learn knowledge from subjective data continually on an ongoing basis to be considered *self aware*.

2.3.2 Reference architecture

Chandra et al. introduces their own new Reference Architecture for designing and describing Computational Self-Awareness and Self-Expression [4]. This abstract and conceptual framework has several advantages. Firstly, it provides a detailed and fine-grained separation of different self awareness and self-expression processes — and hence makes design- and engineering-questions easier and more methodical. Secondly, being a common language, it paves the way for identifying common architectural patterns used to design Comp. SA- & SE-systems. Compared to other Self-Adaptive systems, it also includes the notions of Self-Awareness from the beginning (as the psychology literature suggests SA capabilities is more innate than simply an add-on). And due to the general and implementation-agnostic framework, it can be used on a wide variety of agents systems, or agent-collectives systems.

A schematic of this reference architecture is given below:

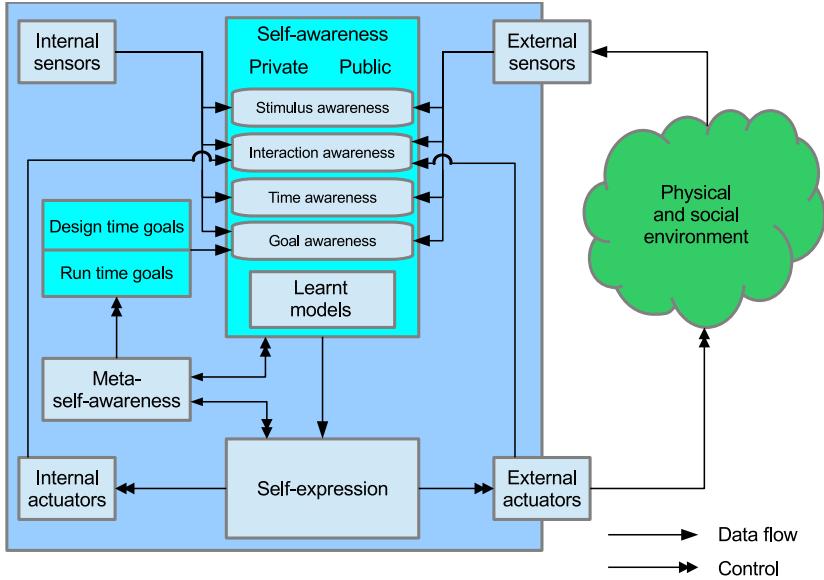


Figure 2.7: Chandra et al.’s [4] reference architecture for a self aware, and self expressive, computing system. Figure credits: Lewis et al. [4].

2.3.3 Conceptual framework

Lewis et al. [13] introduces an inclusive and extensive Conceptual Framework, built and based on the psychological foundation and basis (*as well as expanding on the Reference Architecture*), as introduced earlier in this section.

This conceptual framework is based on the following three high-level concepts and tenets:

1. Levels of self awareness

- **Pre-Reflective self awareness** — ability to subjectively and uniquely perceive and experience phenomena in the physical-/social-environment as well as within itself. A pre-requisite for all other and more advanced forms of computational self awareness, also implicitly introducing the notion of subjectivity.
- **Reflective self awareness** — modelling and conceptualizing various knowledge-/awareness-aspects about an *object*, where the object is what is modelled, or the sensory observation of it at least.
- **Meta-Reflective self awareness** — **Reflective self awareness** processes about its own self awareness processes. Reasoning and modelling of the oneself’s own (as well as potentially other systems’) self awareness properties and processes. Typically its *objects* are 1) the self awareness (modelling) processes themselves, and 2) the output of these processes, i.e. the models themselves. And often the combination (e.g. the memory usage of the learning of an AI model, versus the accuracy of the model — hence potentially calling for a change in architecture or algorithm-choice in a self awareness process) are useful.

2. Aspects of Reflective and Meta-reflective Self-awareness

- All the different aspects and facets of models being learnt and reflected with (like temporal aspects, interactive/causing, identity and state aspects, behavior awareness, appearance awareness, goal awareness, belief awareness capturing uncertainty and trust-levels, expectation awareness combining belief- and time-awareness).
- Not a comprehensive and large list of requirements a system has to check off in order to be considered Self-Aware (due to specific context or system design situations calling for different needs) — but rather a good starting point for considering the aspects captured in reflective and meta-reflective self awareness.

3. Domain of Self-awareness, being the combination of the two;

- Span (subject(s) of reflective self awareness, comprising all entities that contribute to the formation of self awareness)
- Scope (object(s) of self awareness, comprising all the entities observed by the subject's self awareness). E.g., in the case of robotics, the number or amount of neighbouring robots they are aware of and communicate with.

From [13]:

"Using the intersection of all three tenets, we may produce descriptions of the specific self awareness of a system. For example, we might construct a sentence as follows: Peter (span) is aware of Ada's (scope) goal (aspect) to reduce the power usage (object). Further, Ada (span) is aware of her own reasoning (meta-self awareness) about what to do (act) about it."

2.3.4 Challenges & Limitations

Physical constraints like computational processing power, as well as time, limited attention e.g. are all significant constraints which designers of Self-Aware and Self-Expressive Computing Systems have to keep in mind.

2.4 Taking inspiration from natural phenomena

The intriguing, diverse, and complex phenomena of nature have for long served as exciting inspirations to human engineers and researchers [ant-colonies, boids, swarms, beeclust].

Often times, scientists have drawn inspiration from various scientific fields — particularly different fields from ones own — into their own field, for various reasons. Indeed, the translated concepts (from the one domain to the other) will most likely be accompanied with brand new ways to think about ones own domain, as well as other domains again interacting with it (hence having a real opportunity to start a "domino"-like chain-reaction of new ways to think about things emerging). These new ways to think about things (often in ones own domain) — apart from being interesting and intriguing — might be useful, both for ones own field but also for other fields again (especially if thinking long term). For example in the Multi-Agent Systems (MAS) field, it has been

a common practice to study complex biological systems in nature, in order to translate these mechanisms into the technology- and engineering-domain. Such bio-inspired algorithms have been some of the most widely used optimization algorithms throughout history, with e.g. the original paper by M. Dorigo et al. for the ant colony optimization algorithm [6], as of May 2022, having 14 743 citations on Google Scholar. Another example is how D. Goldman et al. [8] took inspiration from e.g. zebra-tailed lizards able to run up to 5 meters per second (normally in desert sand) when designing a robot traversing one of the most challenging terrains to traverse through, granular surfaces like sand that is. Natural phenomena have then been subject to considerable study and modelling—and have in fact created entire scientific fields in which principles from various science- & engineering-dicipline are applied to the physical systems and machines having functions that mimic biological processes, as well as biological processes serving as great sources of inspiration for the engineered systems; *Biomimetics*⁴ and *Bionics*, that is.

One branch of these types of natural phenomena observed and studied are the biological networks of pulse-coupled oscillators [20, 24]. An example of such biological pulse-coupled oscillators as found in nature, particularly of interest for this thesis, is the synchronously flashing fireflies, as e.g. can be seen in a dark forest in Figure 1.1 in Chapter 1. This phenomenon has inspired scientists like e.g. Mirollo & Strogatz [15] and in later time Kristian Nymoen et al. [22], to model and replicate this natural phenomenon in human-engineered systems. Given the periodic and repeating nature of the flashing/firing of the fireflies, modelling a firefly has been done by looking at each firefly as a periodic signal or oscillator. This work [15, 22] then ties into the broader work on synchronizing oscillators which has been subject to study for some time now [1]. What separates Mirollo-Strogatz and K. Nymoen’s approaches from these other and previous oscillator-synchronizing methods, is mainly that here the oscillators are *pulse-coupled* (which the fireflies also can be said to be), as opposed to the more “standard” and constraining *phase-coupled* oscillators. Additionally, K. Nymoen’s approach accounts for synchronizing initially heterogenous frequencies as well.

2.5 Oscillators and oscillatorsynchronization

Oscillators have been used to implement and model a plethora of systems, also biological ones, ranging from designing the locomotion-patterns of swimming robot-amphibia through central pattern generators [10], ..., and as we have already established, modelling synchronously flashing or firing fireflies.

Most importantly for this thesis, an oscillator mainly consists of phase value ϕ and a frequency value ω , much like in a periodical signal like e.g. $\sin(\omega \cdot t + \phi)$. These particular aspects of oscillators will be explained further now.

2.5.1 Phase and frequency

Much of the terminology from [22] is used here. An oscillator i is characterized by its *phase* $\phi_i(t)$, which is—at the start of its periodic signal period—initialized

⁴*Nature*, <https://www.nature.com/subjects/biomimetics> (accessed 2022.05.17)

to 0 and evolves towards 1 at a rate which is called the *frequency* of the oscillator. So, in mathematical terms, the frequency $\omega_i(t)$ is then given by:

$$\omega_i(t) = \frac{d\phi_i(t)}{dt}. \quad (2.2)$$

When oscillator i 's phase is equal to 1 (i.e. when $\phi_i(t) = 1$, or when the periodic signal period is over), we say oscillator i has *phase-climaxed*.

An oscillator-period T is defined as the inverse of the oscillator-frequency ω . In mathematical terms:

$$T = 1/\omega. \quad (2.3)$$

2.5.2 Phase synchronization

2.5.2.1 Miroollo-Strogatz's "standard" phase adjustment

One approach having been used to achieve this in the past is Miroollo-Strogatz's "Standard" phase adjustment in oscillators [15], as seen a sketch of in Figure 2.8.

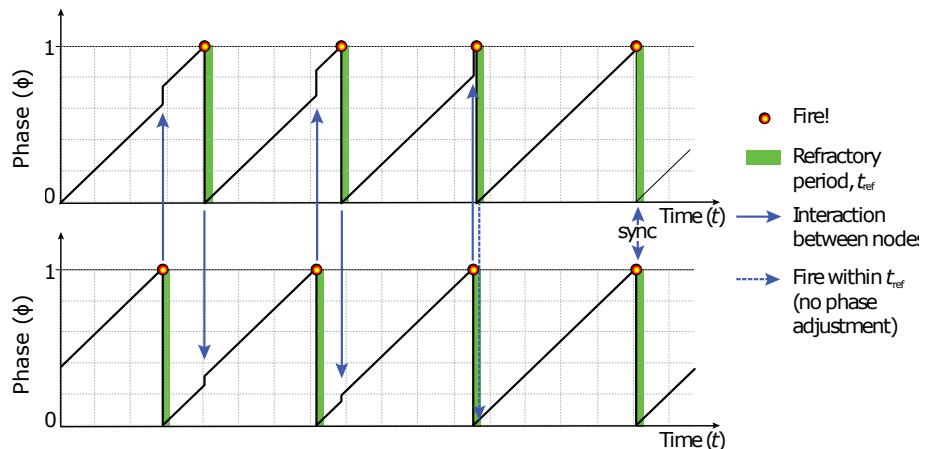


Figure 2.8: A sketch of Miroollo & Strogatz's "standard" phase adjustment. Figure credits: K. Nymoen et al. [22].

Each oscillator gets a new phase, $\phi(t^+) = P(\phi(t))$, according to the **phase update function** (2.4) upon perceiving a "fire"-event from one of the other musical nodes:

$$\phi(t^+) = P(\phi(t)) = (1 + \alpha)\phi(t), \quad (2.4)$$

where α is the pulse coupling constant, denoting the strength between nodes [22], t^+ denotes the time-step immediately after a "fire"-event is heard, and $\phi(t)$ is the old frequency of the oscillator at time t .

So, if e.g. $\alpha = 0.1$, then a musical oscillator's new and updated phase, immediately after hearing a "fire"-signal from another oscillator, will be equal to $\phi(t^+) = P(\phi(t)) = (1 + 0.1)\phi(t) = 1.1\phi(t)$. 110% of its old phase $\phi(t)$, that is.

Hence, and in this way, the oscillator would be “pushed” to fire sooner than it would otherwise (as nodes fire once they have reached phase-climax $\phi(t) = 1$).

Chapter 3

Baseline

In this Active Music System, as shown in Figure 3.1, decentralized harmonic synchronization of a dynamic and arbitrary number of musical nodes (modelled as pulse-coupled oscillators) — involving indirect coordination through stigmergy, as well as self-awareness through knowledge regarding itself internally and in publicly in the context of other nodes — is achieved.

This is done by endowing oscillators with self awareness capabilities amongst other measures, and is achieved by nodes communicating solely through audio signals. The platform Nymoen et al. achieved this on iOS devices shown in Figure 3.1, as each firefly on their iOS devices initially start flashing unsynchronized and randomly, but then eventually manage to start firing synchronously¹.



Figure 3.1: Decentralized oscillators modelled as fireflies achieving harmonic synchrony on iOS devices, after starting flashing unsynchronously. Image: K. Nymoen et al. [22].

As Nymoen designed update functions for synchronizing both phases and frequencies, not only were their firefly-inspired oscillators endowed with self awareness capabilities; they were also designed to adjust their phases and frequencies with (close to) 0 change half-way through the oscillator cycle. These (close to) 0 changes half-way through oscillators's cycles enable individual musical agents to play with different measures (e.g. half or 1/4th as fast) compared to other individuals—being characteristic for musical interaction with several musical individuals—as well as enabling the possibility of achieving the target state of harmonic synchrony (cf. 3.1). Their contribution also differs from previous approaches in that Nymoen et al.'s fireflies only need to transmit “fire”

¹See video at <https://vimeo.com/67205605> (accessed 2022.05.17)

signals every other phase climax, not every single phase climax; in other words, synchronization is achieved despite less communication.

First, the relatively new system target state of *harmonic synchrony* will be expounded and explained in 3.1. Then, the aforementioned update functions for both phase (ϕ) and frequency (ω), explaining the incorporated level of self awareness as well as how a decentralized collective of pulse coupled oscillators can achieve harmonic synchrony, are thus presented in 3.2 and 3.3.

3.1 System target state: harmonic synchrony

The state of harmonic synchrony is defined [22] as the state in which all agents in the musical collective “fire” or “flash”, as described in Subsection 4.1.3, at an even and underlying interval or pulse, a certain number of times in a row. This is not to say all agents will have to “fire”/“flash” simultaneously, as has traditionally been the case for strict synchronization in pulse-coupled oscillators []. How legal and harmonically synchronized oscillator frequencies and phases look like will be described in this Section.

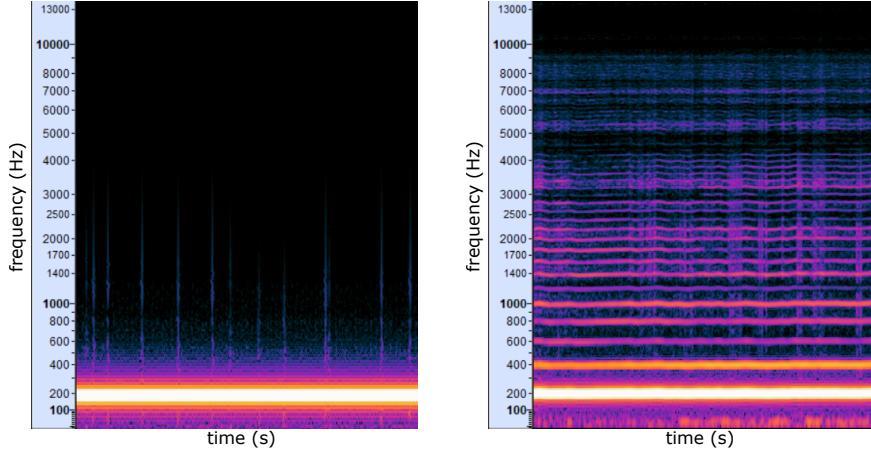
As one is designing and creating an interactive music technology system, one might want to encourage and allow for the playing of various musical instruments at various rhythms/paces, as it might be quite boring if all instruments were played at the exact same measure or pulse. As K. Nymoen et al. [22] reason while discussing their own interactive “Firefly” music-system as well as coining the term of *harmonic synchrony*:

Temporal components in music tend to appear in an integer-ratio relation to each other (e.g., beats, measures, phrases, or quarter notes, 8ths, 16ths).

and

Being an interactive music system, people may want their device to synchronize with different subdivisions of a measure (e.g. some play quarter notes while others play 8ths).

Accommodating for these aspects then, K. Nymoen et al. took inspiration for coining a novel state of synchronization for a decentralized system. This novel state arose from the concept of *harmonics* in the frequency spectrum of a waveform; in that each harmonic wave or overtone has a frequency with an integer relationship to the fundamental (smallest) frequency in a waveform. Such a phenomenon, used as inspiration by Nymoen et al. but not as the basis for any definitions directly, can e.g. be seen in the frequency spectrogram of a humanly hummed G3-tone, depicted in Figure 3.2b. In said example, one can observe the presence of harmonics and overtones having frequencies with integer relationships to the fundamental (smallest) frequency at around 196 Hz, as we see the overtones e.g. has frequencies of around 400Hz, 800Hz, and 1600Hz. We will soon see how such relationships in frequency are used as inspiration for a new definition of a synchronous state; however, the example in Figure 3.2 is only meant to give an idea of where this new state of synchronization stems from, and not as a visual representation of said novel synchrony state.



- (a) The frequency spectrogram of the audible waveform being a monotone and purely generated²G3-tone at 195.99 Hz.
- (b) The frequency spectrogram of the audible waveform being a more-or-less monotone but non-pure G3-tone, hummed and recorded³ by me, as I tried to repeat the tone in 3.2a with my voice.

Figure 3.2: Frequency spectrograms of two different-sounding waveforms of the same G3-tone at 195.99 Hz. Note the absence and presence of harmonics and overtones in waveform 3.2a and 3.2b respectively, as well as the integer-relationships between the fundamental (lowest) frequency and the harmonics in 3.2b. Frequencies in a harmonically synchronized oscillator collective will for the first phase synchronization problem resemble the frequencies in 3.2a, where all frequencies are equal and constant. Conversely, when frequencies can initially be heterogeneous and unequal, as in the phase *and* frequency synchronization problem, the frequencies in a harmonically synchronized oscillator collective will rather resemble, although not completely correspond to, the frequencies in 3.2b, where higher frequencies with integer relationships to the fundamental and lowest frequency can be present at the same time.

More accurately then and inspired by integer relationships in waveforms, although not exactly analogous to the example above in Figure 3.2, Nymoen et al. describe *legal* frequencies the oscillator frequencies has to lie within to be considered *harmonically synchronized*. Frequencies have to lie within

$$\omega_0 \cdot 2^{\mathbb{N}_0}, \quad (3.1)$$

where ω_0 represents the smallest oscillator frequency in the oscillator collective, and \mathbb{N}_0 are the natural numbers including the number zero.

Note that the only difference between Nymoen's definition of *legal oscillator frequencies* and the spectrogram example depicted in Figure 3.2 lies in that for the legal frequencies as defined in (3.1), higher frequencies can only have doubling integer relationships with the lowest fundamental frequency (e.g. twice

²<https://www.szynalski.com/tone-generator/> (accessed 2022.05.17)

³https://github.com/theRealSherapat/CompSA/blob/main/CollectedData/Audio/G3_196Hz_humming.wav (accessed 2022.05.22)

or four times as high), whereas the frequencies of the harmonics seen in the humanly hummed G3 tone (3.2b) additionally had all other positive integer relationships (e.g. three or five times as high) with the fundamental frequency.

3.1.1 Detecting harmonic synchrony

In order to test and evaluate synchronization performance in their firefly-inspired oscillator-system, K. Nymoen et al. [22] develop a measurement used to detect when the system has reached a state of synchrony. Using the firings of the fireflies, some key and well-defined conditions have to be met in order for the fireflies to be deemed *harmonically synchronized*:

Harmonic synchronizaton conditions:

Condition 5): Firing may only happen within a short time period t_f .

Condition 6): Between each t_f , a period t_q without fire events must be equally long k times in a row.

Condition 7): All nodes must have fired at least once during the evaluation period.

Visually, these requirements are shown in Nymoen's sketch for detecting harmonic synchrony in Figure 3.3.

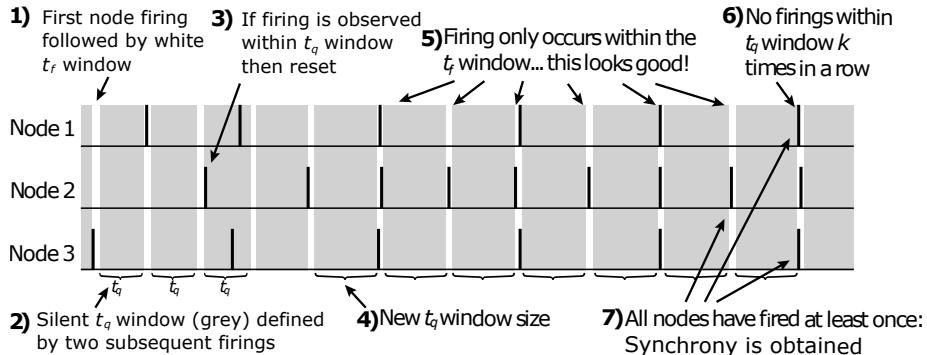


Figure 3.3: Requirements needing to be fulfilled in order to have achieved the target state of harmonic synchrony in an oscillator (here meant by “node”) collective. Image: K. Nymoen et al. [22].

Steps 1) - 4) are means to achieve 5) - 7).

In summary, by utilizing transmitted firings or pulses from the robots in our robot collective, the aforementioned harmonic synchrony conditions can be enforced and checked throughout the synchronization process, in order to detect if, and in that case when, the oscillator network becomes harmonically synchronized. The time having passed from the start of the synchronization-process until the detection of harmonic synchrony will then be defined as the performance score, indicating how fast or slow the oscillators are at synchronizing.

These harmonic synchrony requirements, as well as Nymoen's schema in Figure 3.3, thus constitutes a blueprint for the design of a performance or synchrony measure able to detect the achievement of harmonic synchrony in a decentralized network of pulse coupled oscillators "firing" pulses in order to synchronize to each other; which is exactly what's implemented in this thesis's synchronization system to be explained further in Chapter 4.

3.2 Phase synchronization

In their firefly system, synchronization of the phase values ϕ , representing where the oscillator are in their cycles (e.g. in the start or end of its period), is achieved. By utilizing bi-directional phase shifts, oscillators depending on being in the first or second half of its phase cycle, i.e. $\phi < 0.5$ or $\phi > 0.5$, adjust their phases by either subtracting from the phase, or adding to it. This they do based on their phase update function, $P(\phi)$, as shown in (3.2).

3.2.1 Problem statement

In this first and more simple synchronization problem, which we call the **phase (ϕ) synchronization problem**, we assume homogenous and already-synchronized frequencies $\omega_i = 1\text{Hz}$ for all oscillators i . The oscillator collective hence only has to synchronize phases ϕ_i for all oscillators i , which are initially uniformly random phase values $\in [0, 1]$. Hence, phase adjustments are needed in order to synchronize phases in the pulse coupled oscillator collective.

3.2.2 Synchronizing phases via phase adjustment

K. Nymoen's bi-directional phase adjustment

Apart from altering a similar phase adjustment function to the likes of Miroollo-Strogatz's phase adjustment function (2.4) sketched in Figure 2.8, K. Nymoen et al. [22] introduces an example of a bi directional phase adjustment method. This they do mainly so that their phase adjustment (Adj_ϕ) method would cooperate well with their system's frequency adjustment, in order to achieving their target goal of harmonic synchrony. The difference between a bi directional method of phase synchronization and a one directional one is that using bi directional phase adjustment, oscillators are simply adjusting each other's phases in an excitatory way; they only "push" other oscillators's phases further or higher (positive updates) when firing themselves, never "holding" or "dragging" them back (negative updates).

Hence, this newer approach to phase adjustment works very similarly to the phase adjustment performed in the "standard" **Miroollo-Strogatz** approach presented earlier; the only difference being that now, oscillators update their phases with the slightly more complex **phase update function** (3.2) when hearing a "fire"-event from one of the other musical nodes — allowing for both larger, but also smaller, updated phases compared to the old phases:

$$\phi(t^+) = P(\phi(t)) = \phi(t) - \alpha \cdot \sin(2\pi\phi(t)) \cdot |\sin(2\pi\phi(t))|, \quad (3.2)$$

where t^+ denotes the time-step immediately after a “fire”-event is heard, and $\phi(t)$ is the old frequency of the oscillator at time t .

The fact that new and updated phases can both be larger, but also smaller, compared to the old phases, is exactly what’s meant by the phase-adjustment being ***bi-directional***, or as the authors call it in the paper as using both excitatory and inhibitory phase couplings between oscillators [22].

The effects then of adjusting phases—upon hearing “fire”-events, according to this newest update-function (3.2)—are that the nodes’s updated phases $\phi(t^+)$, compared to their old phases $\phi(t)$, now get decreased if $\phi(t)$ is lower than 0.5, increased if $\phi(t)$ is higher than 0.5, and neither—at least almost—if the phases are close to 0.5. This is due to the negative and positive sign of the sinewave-component in Equation (3.2), as well as the last attenuating factor in it of $|sin(2\pi\phi)| \approx |sin(2\pi\frac{1}{2})| = |sin(\pi)| = |0| = 0$, then if we have $\phi(t) \approx 0.5 = \frac{1}{2}$.

3.3 Phase and frequency synchronization

Nymoen also achieved synchronizing the frequency values ω of oscillators, being the speed at which the oscillator cycles through its cycle or oscillator period at (cf. (2.2)).

This they managed to do by as mentioned and amongs other measures implementing an aspect of *self awareness* in terms of a *synchrony self assessment* $s(n)$ given “fire” event n , representing how much or out of sync an individual robot is.

It will then, having taken into account some middle steps as explained below, use this self assessed sync score $s(n)$ to obtain a frequency adjustment contribution $H(n)$ for “fire” event n , which will tell in what direction and how much the frequency should be adjusted. As many such $H(n)$ values will be calculated and accumulated in a firefly as said firefly hears fire events throughout its oscillator period / cycle—before they are used for frequency updating at the end of the cycle. This way of performing frequency updates are inspired by Werner-Allen’s reachback firefly algorithm (RFA), simply meaning updates are postponed until the end of the cycle instead of happening immediately as a fire event is received.

Now the system used for comparison later on in the synchronization simulator is explained further in detail.

3.3.1 Problem statement

In this second and more challenging synchronization problem, which we call the **phase and frequency ($\phi\&\omega$) synchronization problem**, we no longer assume homogenous and already-synchronized oscillator frequencies like in the first phase (ϕ) synchronization problem when all frequencies were fixed to 1Hz. Now, the oscillators’s frequencies are initialized to uniformly random frequencies ω_i within a certain minimum initial frequency ω_{min}^{init} (e.g. 0.5Hz) and maximum initial frequency ω_{max}^{init} (e.g. 4Hz). Hence, the problem is now more complex; the oscillator collective not only has to synchronize all initially random phases ϕ_i as above, but simultaneously, also all initially random frequencies ω_i for all oscillators i as well. Hence, both phase *and* frequency adjustments are needed in order to synchronize both phases ϕ and frequencies ω in the oscillator collective.

3.3.2 Synchronizing frequencies via frequency adjustment

Only one frequency adjustment / update function was used as a starting point for this thesis and eventually implemented in the Unity synchrony simulator. This one frequency synchronization method, as invented by Nymoen et al. [22], is now presented.

K. Nymoen's frequency adjustment

This approach to frequency adjustment in pulse coupled oscillators, in order to achieve synchronized frequencies, stands in contrast to previous approaches of synchronization where oscillators frequencies are either equal and fixed, oscillators have to send out a pulse every oscillator cycle, or they do not include any explicit self awareness capabilities. This is not the case for Nymoen's frequency adjustment.

In order to achieve *harmonically synchronized* and “legal” oscillator frequencies (cf. 3.1) in the oscillator collective, not only the phases (ϕ) have to be adjusted and eventually synchronized; frequencies (ω) also have to be synchronized through frequency adjustment.

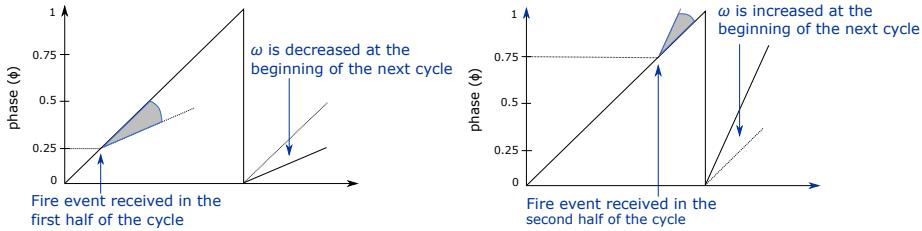


Figure 3.4: An illustration of how frequencies are adjusted after each phase climax (i.e. $\phi = 1$) using Nymoen et al.’s frequency update function. Reproduced from [22].

When it comes to the temporality and timing of when these update functions are used and applied; Musical agents’s phases get updated/adjusted immediately as “fire” events are perceived, whereas agents’s frequencies do not get updated until the end of their oscillator-cycle (i.e. when having a phase-climax $\phi(t) = 1$). This is also the reason why frequencies are updated discretely, not continuously. So-called H-values however, being “contributions” with which the frequencies are to be updated according to, are immediately calculated and accumulated when agents are perceiving a “fire” event — and then finally used for frequency-adjustment/-updating at phase-climaxes.

Each oscillator i updates their frequency, on their own phase-climax (i.e. when $\phi_i(t) = 1$), according to the frequency-update function $\omega_i(t^+)$:

$$\omega_i(t^+) = \omega_i(t) \cdot 2^{F(n)}, \quad (3.3)$$

where t^+ denotes the time-step immediately after phase-climax, $\omega_i(t)$ is the old frequency of the oscillator at time t , and $F(n) \in [-1, 1]$ is a quantity denoting how much and in which direction an oscillator should update its frequency after having received its n th “fire”-signal.

The following steps is how the aforementioned $F(n)$ value is found:

Step 1: the “in/out-of synch” error measurement / score, $\epsilon(\phi(t))$ Describing the error measurements at the n-th “fire”-event, we introduce an error measurement function.

The error measurement function value (3.4) is calculated immediately by each oscillator i , having phase $\phi_i(t)$, when a “fire” event signal from another oscillator $j, j \neq i$ is detected by oscillator i at time t .

$$\epsilon(\phi_i(t)) = \sin^2(\pi\phi_i(t)) \quad (3.4)$$

As we can see from this error-function, the error-score is close to 0 when the oscillator phase $\phi_i(t)$ is itself close to 0 or 1 (i.e. the oscillator either just fired/flashed, or is about to fire/flash very soon). This implies that if it was only short time ago since we just fired, or conversely if there is only short time left until we will fire, we are not much in error or *out-of-synch*.

The error-score is the largest when an oscillator perceives a “fire”-signal while being half-way through its own phase (i.e. having phase $\phi(t) = 0.5$). We could also then ask ourselves, does not this go against the main/target goal of the system, being *harmonic synchrony* — if agents are allowed to be “half as fast” as each other? We could imagine a completely “legal” and harmonically synchronous scenario where two agents have half and double the frequency of each other. The oscillator with half the frequency of the faster oscillator would then have phase $\phi(t) = 0.5$ when it would hear the faster oscillator “fire”/“flash” — leading to its Error-score $\epsilon(0.5) = \sin^2(\pi/2) = 1$, which then makes it seem like the slower oscillator is maximally out of synch, when it is actually perfectly and harmonically synchronized. This calls out for an attenuating mechanism in our frequency update function, in order to “cancel out” this contribution so that perfectly harmonically synchronized agents will not be adjusted further despite their high Error-measurement. As we will see below, exactly such an attenuating mechanism is utilized in our frequency-adjustment method.

This error-measurement/-score forms the basis and fundament for the first component of self-awareness, being the *self-assessed synchrony-score* $s(n)$.

Step 2: The first self-awareness component, $s(n)$ This aforementioned self-assessed synchrony-score, $s(n)$, is in fact simply the median of error-scores ϵ .

If we then have a high $s(n)$ -score, it tells us that the median of the k last error-scores is high, or in other words that we have mainly high error-scores — indicating that this oscillator is out of synch. Conversely, if we have a low $s(n)$ -score, indicating mainly low error-scores for the oscillator — then we have an indication that the oscillator is in synch, hence leading to low error scores, and in turn low $s(n)$ -scores.

In other words, each oscillator hence has a way to assess themselves in how much in- or out-of-synch they believe they are compared to the rest of the agents. This is then the first level of self-awareness in the design.

Step 3: frequency update amplitude- & sign-factor, $\rho(n)$ Describing the amplitude and sign of the frequency-modification of the n-th “fire-event”

received. It is used to say something about in which direction, and in how much, the frequency should be adjusted.

$$\rho(\phi) = -\sin(2\pi\phi(t)) \in [-1, 1] \quad (3.5)$$

For example, if an oscillator i has phase $\phi_i(t) = 1/4$, it gets a value $\rho(1/4) = -\sin(\pi/2) = -1$; meaning, the oscillator's frequency should be decreased (with the highest amplitude actually) in order to "slow down" to wait for the other nodes. Conversely, if an oscillator j has phase $\phi_j(t) = 3/4$, it gets a value $\rho(3/4) = -\sin(3/2\pi) = -(-1) = 1$; meaning, the oscillator's frequency should be increased (with the highest amplitude) in order to get "pushed forward" to catch up with the other nodes.

Acts as an attenuating factor, when $\phi(t) \approx 0.5$, in the making of the H-value — supporting the goal of *harmonic synchrony*.

Step 4: the H-value, and the H(n)-list The following value, acting as "frequency update contributions", is then as previously mentioned calculated immediately after the oscillator perceives another oscillator's "flashing" signal:

$$H = \rho \cdot s \quad (3.6)$$

Here we then multiply the factor $\rho(n)$ representing how much, as well as in which direction, the oscillator should adjust its frequency, together with a factor $s \in [0, 1]$ of the adjusting oscillator's self-assessed synch-score. This implies that $H \in [0, \rho(n)]$. We hence see that the smallest value $H(n)$ can take for the n th "fire"-event is -1, which it does when $\phi = 0.25$ and $s = 1$. The highest value it can take is 1, which it does when $\phi = 0.75$ and $s = 1$. We can also see that even though the self-assessed synch-score $s(n)$ (i.e. the median of error-scores) is high and even the maximum value of 1, thus indicating consistent high error-scores (judging by error-function (3.4)) — the "frequency-update-contribution" $H(n)$ can in the end be cancelled out, as alluded to before, if in fact the amplitude- & sign-factor $\rho(n)$ is equal to 0. Hence, if we have two agents then where the one is twice as fast as the other, and we accept the $H(n)$ -value as the "frequency-update-contribution", the slower oscillator which will hear "fire"-events consistently when it has its phase $\phi \approx 0.5$ (if the agents are synchronized) will, even though it gets a high *out-of-synch score* $s \approx 1$, not "be told" to adjust its frequency more by getting a large "frequency-update-contribution", but in fact "be told" not to adjust its frequency more due to the small or cancelled-out "frequency-update-contribution."

To recall, the self-assessed synch-score $s(n)$ tells an adjusting oscillator how in- or out-of-synch it was during the last m perceived "fire" events — where $s(n) = 0$ signifies a mean of 0 in error-scores, and $s(n) = 1$ signifies a mean of 1 in error-scores. So then if this H -value is to be used to adjust the nodes's frequencies with, the frequency will then be adjusted in a certain direction and amount (specified by $\rho(n)$) — given that the oscillator is *enough* "out of synch"/"unsynchronized" (in the case $s(n)$ is considerably larger than 0).

The H-value says something about how much out of phase the oscillator was at the time the oscillator's n th "flashing"-signal was perceived (and then followingly how much it should be adjusted, as well as in which direction after having been multiplied together with a sign-factor $\rho(n)$), given then that this

H-value also consists of the *self-assessed synch score* $s(n)$ — which again simply was the median of error-scores.

We could look at this H -value as representing the direction and amplitude of the frequency adjustment weighted by the need to adjust (due to being out of synch) at the time of hearing “fire” event n . Or in other words, this H -value is then the n -th contribution with which we want to adjust our frequency with.

Especially interesting cases are when we have $\phi(n) \approx 0.5 \implies \rho(n) \approx \pm 0$, as well as the last m Error-scores $\epsilon(n)$ being close to 0, also leading to $s(n) \approx 0$. In both of these two cases the entire frequency-adjustment contribution H would be cancelled out, due to harmonic synchronization (legally hearing a “fire” event half-way through ones own phase) in the first case, and due to not being out of synch in the latter (having low Error-Measurements). Cancelling out the frequency adjustment contribution in these cases is then not something bad, but something wanted and something that makes sense. If these H -values then are cancelled out or very small, it is indicative of that nodes are already in *harmonic synchrony*, and hence should not be “adjusted away” from this goal state. On the other side, if these H -values then are different (e.g. closer to -1 and 1), it is indicative of that nodes are not yet in *harmonic synchrony*, and that they hence should be “adjusted closer” to the goal state.

The final step: the frequency update function, $\omega_i(t^+)$ Now, we can pull it all together, for Nymoen et al.’s Frequency Adjustment approach for achieving harmonic synchrony with initially randomized and heterogenous frequencies.

When an oscillator i has a phase-climax ($\phi_i(t) = 1$), it will update/adjust its frequency to the new $\omega_i(t^+)$ accordingly:

$$\omega_i(t^+) = \omega_i(t) \cdot 2^{F(n)}, \quad (3.7)$$

where t^+ denotes the time-step immediately after phase-climax, and $F(n)$ is found by:

$$F(n) = \beta \sum_{x=0}^{y-1} \frac{H(n-x)}{y}, \quad (3.8)$$

where β is the frequency coupling constant, y is the number of heard/received “fire-event”s from the start of the last oscillator period to the end (i.e. the phase-climax, or *now*) — and the rest of the values are as described above.

This $F(n)$ -value then, as we see in Equation (3.8), is a weighted average of all the oscillators’s $H(n)$ -values accumulated throughout the oscillator’s last cycle.

Chapter 4

Implementation and experimental setup

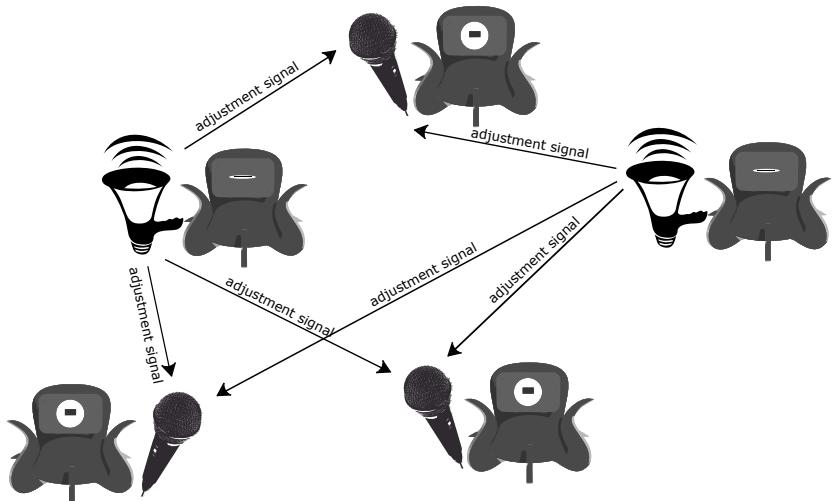


Figure 4.1: Schema: the musical robot collective entraining to synchronize to each other, or more specifically to achieve harmonic synchronization, through performing phase & frequency adjustments. Robots that are not firing at the moment will adjust themselves after hearing a transmitted “fire” / adjustment signal from a neighbouring firing robot.

Envision a decentralized (i.e. no central control) multi-agent collective scenario consisting of musical robots modelled as oscillators. These are solely communicating through brief “fire”-like audio-signals—greatly inspired by Nymanen’s synchronizing “fireflies”. They are not initially synchronized in their firing of audio-signals; but as time goes, they are entraining to synchronize to each other by adjusting their phases and frequencies when/after hearing each other’s audio-/fire signals. If they then, after some time of listening to each other and adjusting themselves accordingly, succeed in becoming synchronized — we then will eventually see “fire”-events/-signals line up at an even underlying

pulse or rhythm. Examples and demonstrations of this process are depicted in Figure 4.1 as well as in Figure 1.2 in Chapter 1.

This chapter gives an overview of the developed musical multi-robot system¹, already existing methods implemented for it, as well as the performance measurement used to evaluate these methods with. The main goal of the implemented system is to allow for individual musical agents in a musical multi-agent collective to interact with each other, in order to achieve emergent and co-ordinating behaviour—in our case synchronization—with varying degrees of self-awareness, collective-sizes, and of difficulty and certainty in the environment and communication. More specifically, the goal with the design is to enable the robot collective to achieve so-called *harmonic synchronization* (cf. 3.1) within a relatively short time. Exactly what is meant by *harmonic synchronization* will be expounded in Section 3.1.

These goals firstly require of the agents the modelling of oscillators with their properties, like phase and frequency, as explained further in Subsection 4.1.2. To allow for interaction and communication between the agents, mechanisms so that the agents can transmit “fire” signals, as well as listen for other agents’s “fire” signals, is necessary as well, and is presented in Subsection 4.1.3.

Firstly, the newly developed synchrony simulator and its components will be presented and expounded in detail. Then, already existing methods implemented in the novel synchronization simulator for achieving the system target goal of *harmonic synchrony* will be described. How the system target state of harmonic synchrony is detected in the synchronization simulator will then be described in Section 4.4.

4.1 Simulator setup: the musical multi-robot collective

Unity² is originally a game development platform, but can also be used to create complex and interactive simulations. Apart from enabling the design, implementation, and testing of various synchronization paradigms and methods—the results of this synchronization can also be clearly and visibly be seen and heard, as showcased in the introductory Figure 1.2 as well as in the video link seen in the corresponding Figure caption.

4.1.1 The simulator environment and its hyperparameters

Hyperparameters further explained

t_{ref}^{dyn} : Roboticists like e.g. K. Konishi and H. Kokame [11] suggest that by reducing the time oscillators in wireless sensor networks are active (i.e. increasing inactive time), power usage can also be reduced—something which makes oscillator systems last longer as well as being better for the environment in the end. Hence, a as large t_{ref} or t_{ref}^{dyn} as possible, seems to have some real benefits in real systems.

¹<https://github.com/theRealSherapat/CompSA> (accessed 2022.05.29)

²Unity Version 2021.2.0f1 is used in the developed simulator.

Notation	Definition
r_i	Individual musical robot i .
R	Musical robot collective consisting of all robots r_i .
ϕ_i	Phase of r_i 's oscillator component.
ω_i	Frequency of r_i 's oscillator component.
$sim\ s$	Simulation seconds. The unit of measurement, w.r.t. to time, used in the thesis and simulation. Simulation time does not always correspond to real / physical time, depending on how fast the hardware is able to run the simulations.

Table 4.1: Notation and terminology used throughout the thesis for system components.

4.1.2 The individual agent: a musical robot

As introduced and presented earlier in 2.2.1.2, our musical robot collective will then consist of models of M. J. Krzyzaniak and RITMO's *Dr. Squiggles*.

The aforementioned (cf. 2.2.1.2) 3D-models of these Dr. Squiggles robots for simulation are reused, with permission, for the simulation system designed in this thesis project. They can e.g. be seen in the introductory Figure 1.2.

Every musical robot or node have certain components, attributes, and characteristics that make it what it is. Such include an oscillator-component, consisting of the agent's oscillator-phase ϕ and oscillator-frequency ω . Notions like "agent", "robot", "firefly", and "oscillator" will be used interchangeably throughout the thesis. The agents have an input-mechanism for hearing/detecting transmitted "fire"-event signals from other agents, as well as an output-mechanism for transmitting or playing such "fire"-/adjust-signals or tones, as is illustrated with the microphone and megaphone respectively in Figure 4.1.

In order to be able to analyse the musical scenario within which they are situated (self-assessment), as well as for adapting their musical output accordingly (self-adaptation), the agents are to some extent endowed with artificial intelligence and self-awareness capabilities. The robots are self-aware of their own phase and frequency, but are unaware of other agents's true phases and frequencies. They also possess the self-assessment capability of evaluating how much in- or out-of-synch they are, as seen in the greater context of the entire robot collective. When the agents hear the transmitted "fire"-/adjust-signals, the agents are intelligent enough to adjust themselves in the direction of the system goal/target state.

Unless otherwise is stated, the heterogenous visual looks of the Dr. Squiggles robots in Unity have no real difference in the simulator and is only that (for visual looks), not implying other values or methods used.

4.1.3 Robot communication: the "fire"-signal

Signals, in various forms, are omnipresent in our world, whether we notice it or not. They are used for guidance: e.g. traffic-lights send visual light-signals to drivers to ensure traffic-flow and collision-avoidance; sirens and alarms fire away with the loudest of audio-signals (sounds) so that its listeners will get out of harms way; our nerves send pain-signals through our nervous-system if we

Hyperparameter	Meaning
$ R $	Musical robot collective size, where R (as defined above) is the set of all musical robots r_i , $i = 1, 2, \dots, R $.
Adj_ϕ	The adjustment method, or update function if you will, used for synchronizing phases ϕ .
Adj_ω	The adjustment method, or update function if you will, used for synchronizing frequencies ω .
t_{ref}	The time (sim s) robots are unresponsive to adjusting themselves (as a consequence of hearing other robots firing) directly after firing themselves.
t_{ref}^{dyn}	Dynamic refractory period. Same as t_{ref} , only that now the refractory period is given as the percentage of robot r_i 's oscillator period (e.g. 10% of $1/4Hz = 0.25s$, i.e. $0.1*0.25s = 25ms$, if the oscillator frequency $\omega_i = 4Hz$).
t_f	Legal firing time during which robots r_i are allowed to fire when being detected harmonic synchrony for. Condition 5) harmonic synchrony detection requirement (cf. 3.1.1 in Chapter 3).
k	Condition 6) harmonic synchrony detection requirement (cf. 3.1.1 in Chapter 3).
t_{max}	Maximum time limit given to musical robot collectives during which the collective might achieve the target state of harmonic synchrony during, unless they are terminated and deemed “synchronization fails.”
ω_{min}^{init}	Minimum initial oscillator frequency a robot's frequency can be assigned at the start of simulations.
ω_{max}^{init}	Maximum initial oscillator frequency a robot's frequency can be assigned at the start of simulations.

Table 4.2: The environmental and collective hyperparameters present and used for the synchronization simulator in Unity.

touch something burning hot to ensure we do not damage our hand severely. If you have ever tried to make a scary sound to scare away an animal—like a cat or crow—they might in fact flee from you if they think your audible signal was scary enough for them. In this specific case, for our musical robots in the musical synchronization simulator, signals will signify calls for adjustment (of oscillator phases ϕ and frequencies ω to be specific).

These aforementioned audio-signals, also referred to as “fire” signals, “flash”-signals, or adjust-signals, are transmitted whenever an agent’s oscillator *peaks* or *climaxes* (i.e. after its cycle or period is finished, having phase $\phi(t) = 1$) — or actually after every second *peak*, as a way (discovered by K. Nymoen et al. [22]) to attain the system target goal of *harmonic synchrony*, to be elaborated upon in Section 3.1.

The “fire” signals are short and impulsive tones that the agents output through their loudspeakers. These short audio-signals/sounds “wildly” transmitted or played into the environment are then the only means of communication within the multi-agent collective, implying that the agents are pulse-coupled, not phase-coupled, oscillators. In other words, our agents will communicate

<i>Hyperparameter</i>	<i>Meaning</i>
α	Phase coupling constant. The larger the α , the more a robot adjusts its phase when hearing “adjustment signals” from neighbouring robots.
Adj_ϕ	Phase adjustment method used to synchronize phase values ϕ with.
β	Frequency coupling constant. The higher the β , the more robots adjust their frequencies ω_i .
Adj_ω	Frequency adjustment method used to synchronize frequency values ω with.
m	Error memory length. The larger the m , the more error scores ϵ robots remember with which they use in their calculation of the self assessed sync score s .
k_s	Integer number of nearest neighbour robots that are within a robot’s self awareness scope, and hence telling how many neighbour robots a robot listens to for fire events.
d_s	Physical radius (in Unity units) around a robot where robots closer to the robot than d_s are within said robot’s self awareness scope, and hence telling how far out from itself a robot listens for neighbours’s fire events.

Table 4.3: The hyperparameters possible to alter for the individual musical robots in the synchronization simulator in Unity.

and co-ordinate with each other through the very typical multi-agent system concept of *stigmergy*.

When an agent detects a “fire”-/adjust-signal, the agent will adjust its own oscillator-properties (phase ϕ and frequency ω), depending on which type of problem the agents are to solve. No individual agent is directly able to adjust or modify the state or properties of any other agent, only its own.

4.1.3.1 Under the hood in the simulator

Three distinct *self awareness scope* (cf. **Domain** in 2.3.3) scenarios are implemented for each individual musical robot in Unity:

1. **k_s nearest neighbours self awareness scope:** Each individual robot hears the k_s nearest neighbouring robot’s “fire” signals (cf. 4.1.3). In this scenario, at least for small k_s values, robots can have more limited and local knowledge.
2. **Radial d_s self awareness scope:** Each individual robot hears neigbouring robots’s “fire” signals within a radius d_s around it. In this scenario, robots’s self awareness scope is more focussed on the spatial locations of other robots, and can also simulate the robot only hearing robots close in space to itself.
3. **Global self awareness scope:** Each individual robot hears all other neighbouring robots’s “fire” signals. In this scenario, robots have maximum

and global knowledge when it comes to awareness of other neighbouring robots.

In this way, the degree to which robots are self aware of or communicating with other robots is increasing. The effects of increasing or decreasing this degree of self awareness are shown in Chapter 5, and especially in 5.1.5 and 5.2.3.

4.1.3.2 Audible to human simulator observer

Fire signal design The fire signal audible to the human observer watching and listening to the musical robots synchronizing to each other should be distinct and short enough so that one can clearly distinguish between which robots are firing and when. At the same time, the firing-sounds should to the human observer not be too sharp and loud to listen to so that it is uncomfortable observing the musical robot collective. Keeping these aspects in mind, some relatively soft and pleasant firing-sounds from an instrument usually associated with harmonious music—the harp—were produced and developed for the synchronization simulator.

Some manually and empirically perceived harmonious (in terms of sounding good when played together) musical tones were digitally reproduced and recorded in the form of harp plucks. This was achieved using the music-making system and digital synthesizer LMMS, and a digital string-instrument in the form of a LMMS-plugin from DSK Music³. The musical tones, perceived to be harmonious when being heard simultaneously and reproduced with the digital harp instrument, were some of the very first tones perceived in the intro of Pogo’s song “Strangerous⁴.“ One possible avenue to explore in order to find harmonious chords or tones—when played together—in a more automatic approach, live and online during simulation, is discussed in Further work in Chapter 6.

However, the harp-sounds produced with DSK’s LMMS-plugin are primarily long and constant harp plucks, as shown in Figure 4.2a. Using such harp plucks as a fire-sound directly can make it hard to distinguish to the human ear when multiple robots are playing these frequently and simultaneously. The harp-sounds thus need to be slightly edited—which they in our implementation were in the audio-editing program Audacity—so that the long and constant harp pluck became more of a “quick” and distinguishable (when played often by several agents) “sound-bullet”—essentially what we want in a solid “fire” signal sound. The “before” and “after” of such an audio-editing process is depicted in Figure 4.2. Edits performed on waveform 4.2a to obtain waveform 4.2b consists of effects like “Fade Out”, to dampen the sound in the tail of the waveform and get a shorter sustain as it is called, and hence obtaining a “quicker” sound. It also includes, in the complete beginning, the “Amplify” effect, in order to get a as high “attack” or maximum amplitude as wanted to make the fire-sound audible enough before it quickly decays. And lastly, obviously it was needed to cut the waveform accordingly, specifically from the time it had amplitude ≈ 0 to the end of the waveform.

³<https://www.dskmusic.com/dsk-world-stringz-updated/> (accessed 2022.05.17)

⁴<https://www.youtube.com/watch?v=cRzcsXDBn8g> (accessed 2022.05.17)

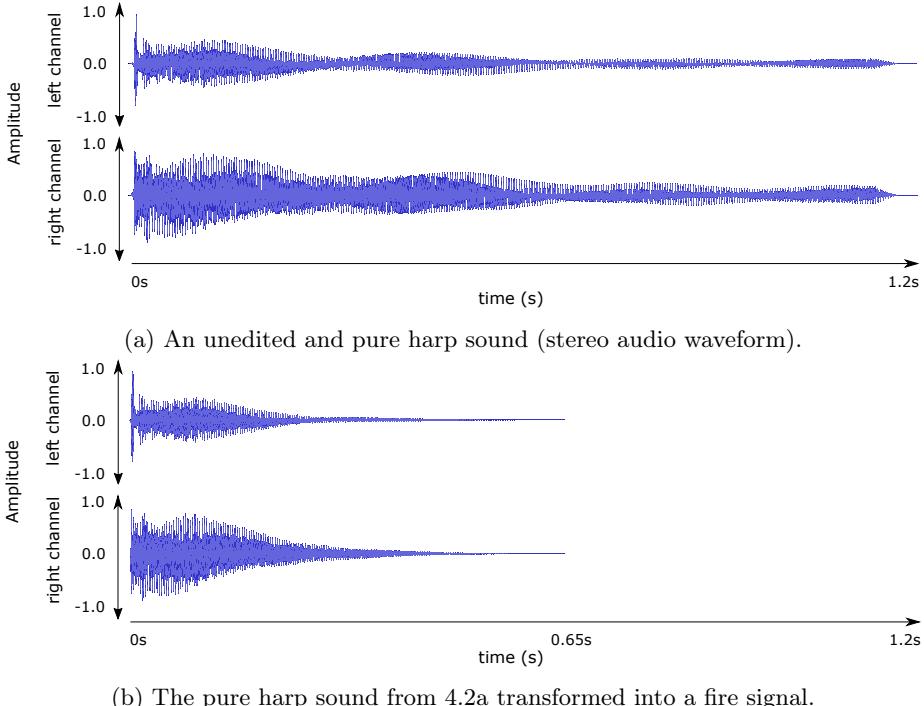


Figure 4.2: Here we see the starting point and end result of a fire sound design process, where the longer and pure harp sound in 4.2a got edited and transformed into the more rapid fire sound in 4.2b.

Frequency based fire-sound assignment Given that our developed system is not simply a synchronization-system, which could have been implemented with simpler electrical circuits e.g. [11], but also a musical robot system — the display and expression of the robots synchronizing and becoming synchronized should have a musical dimension, as well as providing a clear way to see whether the robots are getting synchronized or not. Therefore, a musically interesting and meaningful way of signalizing the synchronization process was needed. The answers to the question of what is different between robots throughout the simulation plays a key part here; namely, e.g. the robot's oscillator frequencies. Different musical tones were recorded (again with a harp instrument) and later edited into “fire” signals, as described above, with various signal lengths given their pitch. The idea was to assign, online and dynamically throughout the simulation run, “fire”-signal-transformed musical tones with higher pitch (or waveform frequencies) to musical robots with the highest oscillator frequencies in the robot-collective. This is done by calculating a frequency percentage $\omega\%$ a given robot r_i has, as in the formula given in 4.1 and visually in Figure 4.3.

$$\omega\% = \frac{\omega_i - \omega_0}{\omega_{max} - \omega_0}, \quad (4.1)$$

where ω_i is the oscillator frequency of robot r_i , ω_0 is the fundamental or smallest oscillator frequency currently in the robot collective, and ω_{max} is the highest oscillator frequency currently in the robot collective.

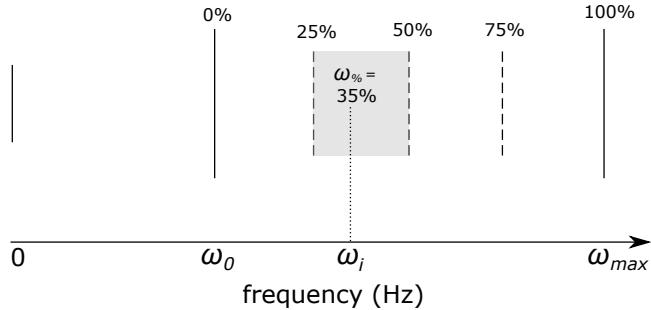


Figure 4.3: Schema of frequency based fire signal assignment in an example where four fire sounds with increasing frequency or pitch exist, and the robot is assigned the second lowest due to its oscillator frequency ω_i .

This percentage is then used in a manner like in the schema given in Figure 4.3, where the frequency percentage $\omega\%$ is decided in which region, corresponding to a lower or higher pitch fire sound, the robot's frequency belongs to. Hence, robots are assigned fire sounds with higher pitch or frequency the higher frequencies the robots themselves have, and similarly fire sounds with lower pitch or frequency the lower oscillator frequencies the robots have.

4.1.3.3 Robot instantiation

Robots are spawned randomly in the Unity scene within a circle from origo until a certain radius according to the empirically found, through trial and error, formula for not too crowded robot collectives in (4.2).

$$radius_{spawn} = \frac{|R|}{6} \cdot r_{width} + 4, \quad (4.2)$$

where $r_{width} = 5.7$ is the measured physical width or diameter in Unity of an individual robot r , and the term 4 simply is to add a bit more space to avoid crowding. The simple model was found through trial and error, and will make the robots spawn kind of like in a round pool of water.

4.2 Synchronizing oscillator phases

If we first assume constant and equal oscillator frequencies in our agents, we can take a look at how the agents adjust their initially random phases in order to synchronize to each other. Examples of how phase synchronization can look are given in Figure 4.4 and 4.5. A more formal requirement and definition for phases to be deemed harmonically synchronized is also given below.

For both of the implemented phase synchronization methods, Miroollo & Strogatz and Nymoen's, the timing of when the update functions are used and applied is the same; Musical agents's phases get updated immediately as "fire" events from neighbouring robots are perceived.

Legal phases definition:

In an oscillator collective with heterogenous oscillator frequencies ω_i , oscillator phases ϕ_i are harmonically synchronized when all oscillators's phase climaxes (i.e. $\phi = 1$) or resets (i.e. $\phi = 0$) occur exactly $T_{min} \cdot N$ seconds later after t_{synced} , where $T_{min} = 1/\omega_{max}$ for the highest oscillator frequency in the collective ω_{max} , t_{synced} is the time after the oscillator collective has become harmonically synchronized and the oscillator with the highest frequency ω_{max} has a phase climax or reset, and N are all natural (positive integer) numbers.

4.2.1 Implementing and verifying Miroollo-Strogatz's phase adjustment

Miroollo-Strogatz's approach for synchronizing phases in oscillators, as introduced in 2.5.2.1, is implemented in the Unity simulator, and each agent is endowed with **phase update function** (2.4) with which they adjust themselves according to when perceiving a "fire"-signal as described above.

The verification that this works in the newly built synchronization simulator was performed by dumping all agents's phase values $\phi(t)$ during simulation runs. A plot of these $\phi(t)$ -values, evolving through simulation time in seconds, is shown in Figure 4.4.

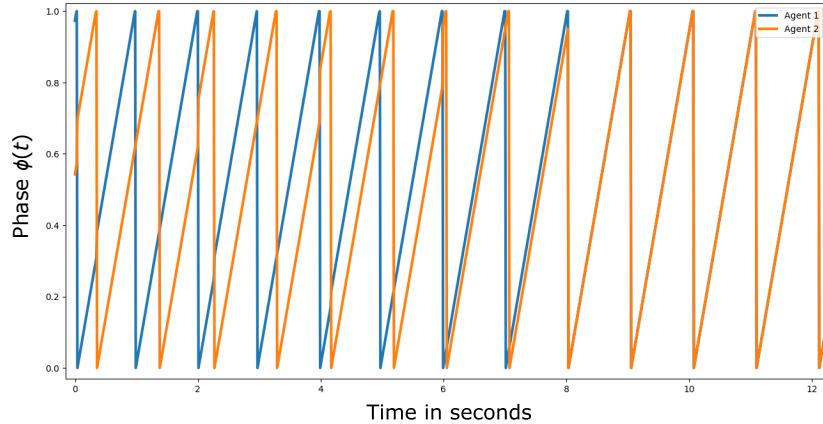


Figure 4.4: "Standard" phase adjustment with Miroollo-Strogatz's approach

4.2.2 Implementing and verifying Nymoen's bi-directional phase adjustment

Nymoen's approach for synchronizing phases in oscillators, as introduced in Section 3.2.2, is implemented in Unity, and each agent is endowed with **phase update function** (3.2) with which they adjust themselves according to when perceiving a "fire"-event as described above.

The verification that this works in the newly set-up simulator-environment was performed by analysing carefully all the agents's phase values $\phi(t)$ throughout a simulation run. Such an analysis plot can be seen in Figure 4.5.

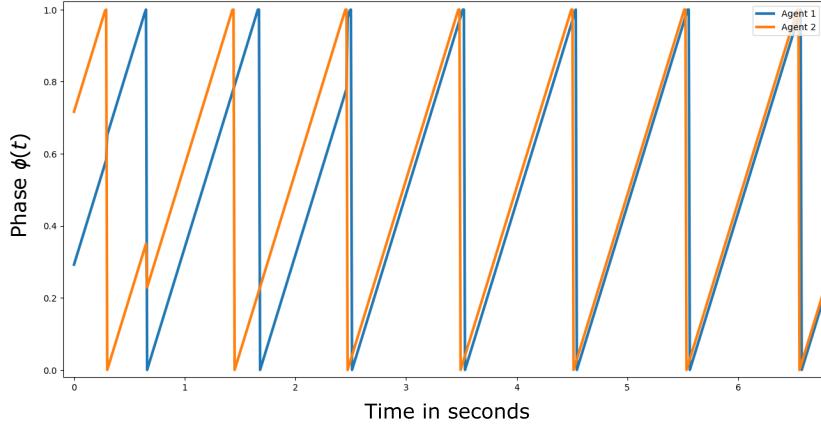


Figure 4.5: Bi-directional phase-adjustment with Nymoen’s approach

4.3 Synchronizing oscillator frequencies

When we open up for the possibility for heterogeneous frequencies in our musical agent collective, we open up to exciting musical aspects like the playing of diverse rhythmic patterns as e.g. mentioned in Section 3.1, but we then also need to not only synchronize phases, but also frequencies, simultaneously. An example of how Nymoen’s frequency synchronization, implemented in the current Unity synchronization simulator, can look like is given in Figure 4.6.

For Nymoen’s frequency synchronization method, the timing for when the update function is to be used and applied is a bit peculiar, and only happens once every phase climax ($\phi = 1$). This is due to its utilization of the so-called reachback firefly algorithm, as explained in Subsection 3.3.2 and displayed in the corresponding Figure 3.4.

We hence now introduce randomly initialized, non-constant, and heterogeneous oscillator frequencies in our musical agents. The agents are now required to synchronize their initially different and random frequencies, so that frequencies are “legal” and *harmonically synchronized*. Such “legal” frequencies are now described clearly in detail.

Legal frequencies definition:

Building upon Nymoen’s definition of harmonic synchronized oscillator frequencies (cf. (3.1)), a concise definition for the legal frequencies can be given as follows. All musical oscillators i , in a harmonically synchronized state, will have frequencies ω_i which are element in the mathematical set

$$\Omega_{legal}(\omega_0) = \omega_0 \cdot 2^{\mathbb{N}_0} = \{\omega_0, 2\omega_0, 4\omega_0, 8\omega_0, \dots\}, \quad (4.3)$$

where ω_0 is the lowest frequency in the oscillator collective (or the fundamental frequency if you will), and \mathbb{N}_0 are the natural numbers including the number zero.

If e.g. the smallest oscillator frequency in the musical oscillator collective

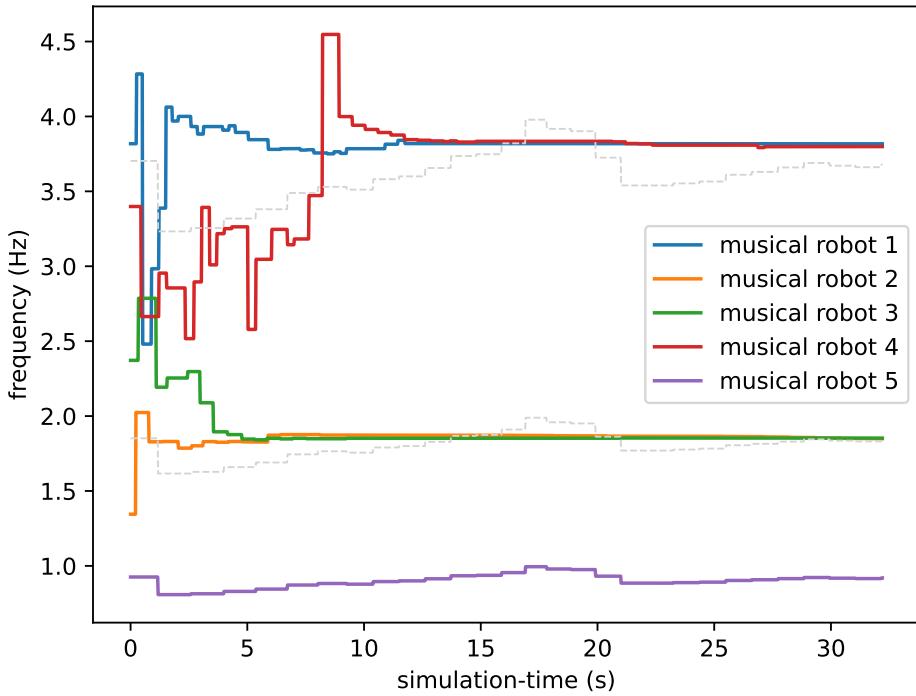


Figure 4.6: Frequency synchronization for five robots in a Unity synchronization simulation run achieving harmonic synchrony after 30 seconds. Note how the legal harmonic frequencies (dashed gray lines) are defined by the lowest—fundamental—frequency ω_0 in the robot collective fluctuating slightly below 1Hz, correctly leading to the legal frequencies right below 2 and 4 Hz.

(ω_0) was equal to 1.5Hz, *legal* frequencies the rest of the oscillators in the collective could have would be $\Omega_{legal}(1.5\text{Hz}) = \{1.5\text{Hz}, 3\text{Hz}, 6\text{Hz}, 12\text{Hz}, \dots\}$.

Hence, in terms of frequencies ω_i for all oscillators i in the oscillator network, we have *harmonically synchronized* and *legal* oscillator frequencies ω_i if (and only if)

$$\omega_i \in \Omega_{legal}(\omega_0), \forall i. \quad (4.4)$$

This state of *harmonic synchrony* is then the system goal state K. Nymoen et al. achieve using their phase and frequency update functions, as explained in Sections 3.2 and 3.3; and is also the target or goal state we want to continue achieving and experimenting for in this thesis.

Implementation details:

In the newly proposed Unity simulator environment, the previously introduced self-assessed sync score $s(n)$ (in 3.3.2), encompassing the approach's main self awareness capabilities, is implemented as the median of a list containing m error-scores ϵ . Such an error score list of length m is easily implemented by a list called *errorBuffer*:

$$errorBuffer = \{\epsilon(n), \epsilon(n-1), \dots, \epsilon(n-m)\}, \quad (4.5)$$

then leading to:

$$\begin{aligned} s(n) &= \text{median}(errorBuffer) \\ &= \text{median}(\{\epsilon(n), \epsilon(n-1), \dots, \epsilon(n-m)\}) \in [0, 1], \end{aligned} \quad (4.6)$$

where n is the latest observed “fire-event”, and m is the number of the last observed “fire”-events we would like to take into account when calculating the self-assessed synch-score.

Regarding the “frequency-update-contributions” (the H -values described in 3.3.2) in my Unity-simulator, all the calculated H-values are accumulated and stored in an initially empty C#-list (of floats), referred to as $H(n)$, at once they are calculated. The $H(n)$ -list is then consecutively “cleared out” or “flushed” when its H-values have been used for the current period’s frequency adjustment (i.e. at the phase climax, when $\phi(t) = 1$), and is then ready to accumulate new H-values during the next period.

By using K. Nymoen et al.’s frequency adjustment method, synchronization of initially random frequencies in our newly developed Unity simulator is achieved.

The details of how the Dr. Squiggles in Unity adjust and eventually synchronize (harmonically) their oscillator frequencies, is shown in Algorithm 1 and 2. Algorithm 1 outlines how the “frequency adjustment contributions” H get calculated as soon as the Dr. Squiggles hear a neighbour Dr. Squiggle sending a “fire” signal; whereas Algorithm 2 shows how these previously calculated $H(n)$

values are used to obtain the new and updated oscillator frequency $\omega_i(t^+)$ with which individual oscillators assign their frequencies to after each phase climax (i.e. when $\phi = 1$).

Algorithm 1: Calculating frequency update contributions H for a robot

```

1: procedure ( $\phi(t), m$ )
2:   if not in refractory period then
3:      $\epsilon \leftarrow \sin(\pi\phi(t))^2$                                 ▷ Step 1
4:   else
5:      $\epsilon \leftarrow 0$ 
6:   end if
7:    $\epsilon$  is added to errorBuffer
8:    $s \leftarrow \text{median}\{\epsilon(n), \epsilon(n-1), \dots, \epsilon(n-m)\}$       ▷ Step 2
9:    $\rho \leftarrow -\sin(2\pi\phi(t))$                                          ▷ Step 3
10:   $H \leftarrow \rho \cdot s$                                               ▷ Step 4
11:   $H$  is added to list  $H(n)$ 
12: end procedure

```

Algorithm 2: Updating robot frequency with calculated H contributions

```

1: procedure ( $H(n), \beta$ )
2:    $cycleHSum \leftarrow 0$                                          ▷ Step 5 started
3:    $y \leftarrow \text{length}(H(n))$ 
4:   for all  $H \in H(n)$  do
5:      $cycleHSum \leftarrow cycleHSum + H$ 
6:   end for
7:    $F(n) \leftarrow 0$ 
8:   if  $y$  not 0 then
9:      $F(n) \leftarrow \beta \cdot cycleHSum/y$                                ▷ Step 5 finished
10:   end if
11:   Emptying  $H(n)$                                               ▷ for new contributions next oscillator cycle
12:    $\omega_i(t^+) \leftarrow \omega_i(t) \cdot 2^{F(n)}$ 
13: end procedure

```

4.4 Detecting harmonic synchrony

The whole detection of harmonic synchrony in all its simplicity boils down to utilizing Unity's concept of Coroutines, meaning a process that can go on and be distributed across multiple frames in the simulator, as well as the firing of the robots as they every second phase climax ($\phi = 1$) transmit a fire signal the simulator picks up on. As this in itself took several hundred lines of code⁵ to implement, the exact details of how it is implemented will not be expounded here.

The performance measurement is used in the synchrony simulator to evaluate and test the multi-robot collective's ability to achieve harmonic synchronization. As mentioned in Subsection 3.1.1, K. Nymoen et al.'s requirements

⁵<https://github.com/theRealSherapat/CompSA> (accessed 2022.05.29)

and illustrations (cf. Figure 3.3) for achieving *harmonic synchrony* serve as a blueprint or guide for how to similarly implement our synchrony or performance measurement. This performance measurement should be able, during synchronization simulation, to detect if harmonic synchronization has been achieved in our decentralized oscillator network. The successful triggering of this detection will then in turn terminate the synchronization simulation run and save to a dataset the time it took to synchronize (the performance score), in the case of a “synchronization success”—an example of which can be seen in Figure ???. In this figure, one can see at which times (sim s) the musical robots fired during a synchrony simulation run ending up in either harmonic synchrony, or as a “synchronization fail.”

The resulting and corresponding performance scores obtained using this performance measurement will then take values of the termination time (sim s) it takes for the robot collective, from the start of the synchronization simulation, to achieve the system target state of *harmonic synchrony*, as specified in Section 3.1.

Now if **Conditions 5)-7)** from Subsection ?? are kept, we have harmonic synchrony.

My specific implementation of the synchrony measurement essentially consists of enforcing all the requirements or rules listed in ??, given some constant t_f - and k -values (e.g. 80ms and 8 respectively [22]). And again—to recall from ??— t_f is the short time-window within which nodes are allowed to fire at each beat, and k represents how many times nodes have to fire at even underlying pulses/beats in a row without changing the t_q -period—before becoming harmonically synchronized.

The requirement of firing evenly k times in a row with identical t_q -periods can be—and in fact is in our implementation—enforced by incrementing an integer variable *towards_k_counter* after a ‘legal’ t_f -window has occurred (i.e. one or more nodes fired inbetween the onset and ending of the t_f -window), and conversely by resetting *towards_k_counter* to 0 when an illegally transmitted firing was heard during a ‘silent’ (or so it was supposed to be at least) t_q -window, hence restarting the synchrony-detection process—as can be seen occurring several times in Figure 4.7.

Note that in the specific simulation run in Figure 4.7, the robots were on their way to achieve harmonic synchrony five times before the 10th second of the synchronization-simulation already, but since one or more of them fired ‘illegally’ (i.e. inside a t_q -window), they were consequently ‘punished’—or rather deemed ‘not synchronized enough yet’—by getting their counter reset to 0. Eventually however, through further phase and frequency synchronization, the multi robot collective was in this case after 12.5 seconds able to achieve harmonic synchrony, when the “even beat” counter became equal to k , as well as all other requirements for achieving *harmonic synchrony* was met. Note that this gives us a sense of how synchronized the robot collective is over time; the more even beats the robots have in a row, the closer they are to achieving harmonic synchrony.

Initially, the t_q -period/-window is not initialized, as it entirely depends on the frequencies to which the robot-collective converges to; however, when an illegal firing (i.e. a firing perceived during a t_q -window) occurs— t_q is also then reset itself to a hopefully more correct value, given by the following formula, which is visually explained further in the schema in Figure ??:

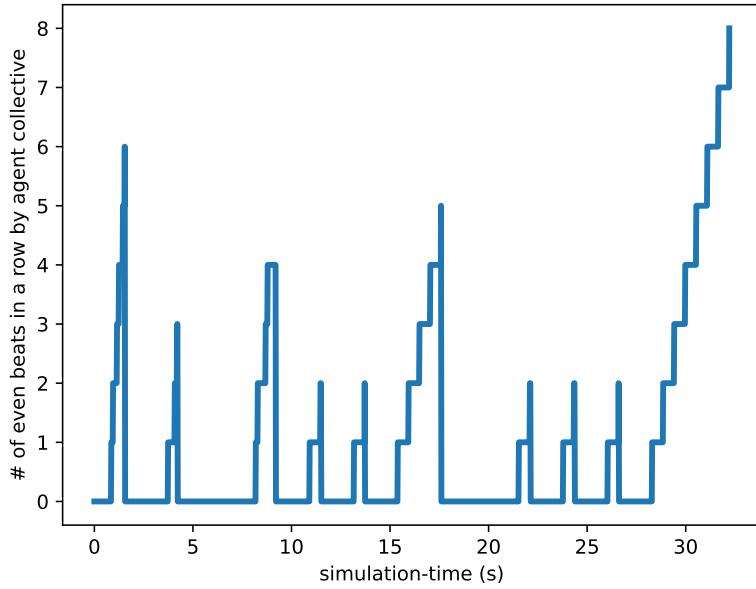


Figure 4.7: A **synchrony evolution plot**, displaying the temporal recording of the *towards_k_counter* variable throughout a synchrony simulation run in Unity. The counter is incremented as the robot collective fires evenly within ‘legal’ t_f windows, and is conversely reset to 0 if illegal firings during ‘silent’ t_q windows are heard.

$$\begin{aligned}
 t_q^* &= \text{late_}t_q^*\text{_defining_timevalue} - \text{early_}t_q^*\text{_defining_timevalue} - t_f \\
 &= \text{median}(t_1, t_2, \dots, t_m) - \text{median}(t_1, t_2, \dots, t_n) - t_f,
 \end{aligned} \tag{4.7}$$

where n fire-events were recorded during the earlier t_f -window, and m fire-events were recorded during the later t_f -window.

If a certain amount of time, e.g. 5 simulation minutes [22], has gone without the detection of harmonic synchrony occurring, the simulation run is terminated as a “fail”.

Chapter 5

Experiments and results

This chapter presents the experiments set up and performed in the novel synchronization simulator in Unity, as presented in Chapter 4, for certain configurations of musical robot collectives. Effects of the individual musical robots's hyperparameters on the collective achievement and performance of achieving harmonic synchrony are presented. Some examples are hyperparameters which determines how much each musical robot will adjust itself after hearing a transmitted fire signal from a neighbouring robot; α for phase adjustment and β for frequency adjustment.

The main performance scores presented in this chapter will consist of synchronization times given in simulation time (s) (i.e. how long it takes robot collectives to reach the state of harmonic synchrony if they ever do), accompanied by the respective and corresponding error rates during the belonging simulation runs (i.e. the percentage of robot collectives out of e.g. 30 runs failing to reach harmonic synchrony before the maximum time limit of e.g. 5 simulation minutes). If a musical robot collective then never reaches the target state of harmonic synchrony within the maximum time limit, this simulation run will be regarded as a "synchronization fail", and we will then not regard its termination time (in simulation time seconds) as *harmonic synchronization time*—but simply that, that simulation run's termination time.

As previously mentioned, relating to the different robot colors like turquoise, red, and green, all robots are homogenous apartly from the visual; all individual robots have for experiments the same hyperparameters unless otherwise is explicitly stated.

5.1 Phase synchronization

This is the section where experiments attempting to synchronize for the first and simpler problem, namely synchronizing only the phases ϕ_i of all agents i , are presented and analyzed resuls for. These are then experiments where all musical robots have an equal and fixed frequency, only adjusting phases, in order to entrain to synchronize their phases to each other until reaching the target state of harmonic synchrony. Or in other words, all experiments in this phase synchronization section have the two hyperparameters shown in Table 5.1 in addition to the individual experiment hyperparameters listed.

Adj_ϕ	Adj_ω
Mirollo & Strogatz	None

Table 5.1: Hyperparameter setup for phase (ϕ) synchronization experiments.

5.1.1 Initial validation

To see that synchronization is actually happening like intended, relevant quantities and simulation values are recorded for a phase synchronization simulation run in Unity and displayed in Figure ??.

5.1.2 Reproducing baseline results

Experiment setup:

In order to see that our developed synchronization simulator in Unity yields more or less the same results as Nymoen’s results, similar experiments as reported in their paper are performed here. These tell us whether differences in performance, in terms of synchronization times (sim s), is simply due to implementation differences, or actually because of the utilized synchronization methods and hyperparameters in question.

First off, Mirollo & Strogatz’s phase adjustment method (as presented in 2.5.2.1) is experimented with for the initial phase (ϕ) synchronization problem, for varying phase coupling constants α . Hyperparameters are set in the simulator as shown in Table 5.2. Results can be seen in Figure 5.1.

$ R $	β	t_{ref}	t_f	k	t_{max}
6	0	50ms	80ms	8	5min

Table 5.2: Hyperparameter setup for ϕ synchronization baseline reproduction experiment.

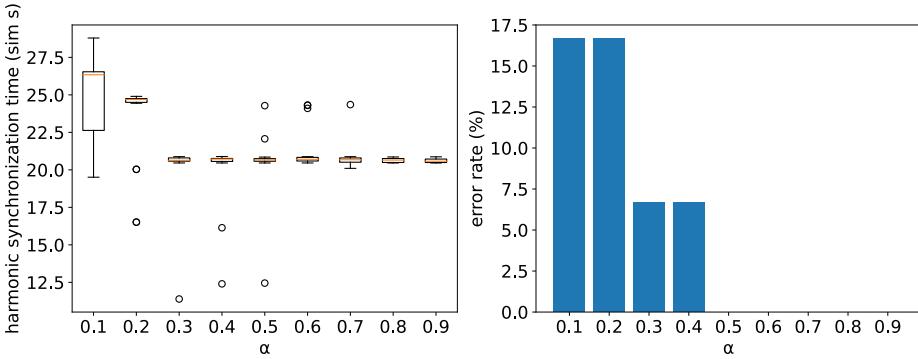


Figure 5.1: Harmonic synchronization times (s) for 6 robots with initially random and unsynchronized phases but equal and fixed frequencies (1Hz), for varying phase coupling constants α . 30 simulation runs per α are reported.

Experiment results:

Experiment analysis:

5.1.3 Hyperparameter tuning

Experiment setup:

Here we tune the hyperparameters α and t_{ref}^{dyn} for several robot collective sizes, according to performance scores of how long it takes robot collectives on average to achieve harmonic synchrony. Exactly these two specific hyperparameters are experimented with mostly since they empirically seem to be the most important ones to set correctly before starting the simulator; that is, in order for the robots to actually manage achieving harmonic synchrony. Additionally, the point of K. Konishi and H. Kokame (cf. 4.1.1) was also remembered.

The specific values of hyperparameters to test synchronization times were chosen based on an initial and identical but failed experiment (not reported) where the tested values for t_{ref}^{dyn} were showing some interesting effects for various robot collective sizes. Simple and limited insight was also collected from trial and error in the complete beginning when trying to facilitate stable and successful simulation runs. And furthermore, after seeing results found in this thesis for α already, we also got some more ideas for which values of α could be interesting to investigate further.

Other constant hyperparameters not being tuned or experimented for in this experiment are shown in Table 5.3. The hyperparameter tuning results, where the effects of tuning hyperparameters α and t_{ref}^{dyn} for various musical robot collective sizes, are shown in Figure 5.2. Again, since we are synchronizing for the phase (ϕ) synchronization problem, now only phases are initially unsynchronized, and frequencies are fixed and constant (1Hz) throughout the simulation runs.

β	t_f	k	t_{max}
0	80ms	8	5min

Table 5.3: Hyperparameter setup for ϕ synchronization hyperparameter tuning experiment.

Experiment results:

As we can see in Figure 5.2, there are generally low standard deviations which make errors in the plots barely visible. We also can notice that the largest differences are seen for lower α values, as average harmonic synchronization times are more and more overlapping and similar for larger α values.

A general pattern we can see is that larger robot collective sizes (larger $|R|$ values) handle lower phase coupling constants α better; in that the music collectives both achieve harmonic synchrony in lower average times than the smaller robot collectives, as well as achieving lower error rates. To this latter point, note that for the lowest α value of 0.001, robot collectives with size $|R| = 3$ and 10 were not able to achieve harmonic synchrony at all during any of the simulation runs (i.e. having 100% error rate), whereas larger collective sizes like $|R| = 25$ and $|R| = 50$ managed to achieve harmonic synchrony despite the low α value.

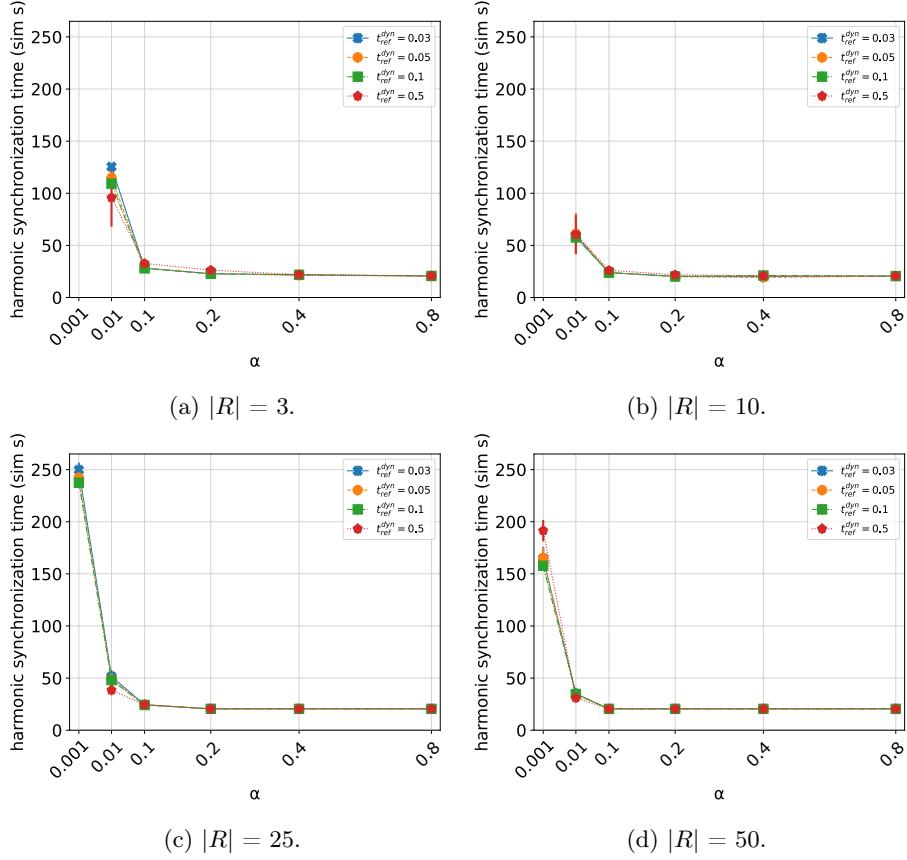


Figure 5.2: Average harmonic synchronization times are plotted in errorbar plots, where the standard deviation is the error. 30 simulation runs per α and t_{ref}^{dyn} pair are reported, unless simulation runs ended up as “synchronization fails.”

Experiment analysis:

The explanation for this latter observation regarding the seemingly increasing robustness with increasing collective sizes $|R|$ can very well lie in the fact that, as we remember, the phase coupling constant α represents how much each robot will adjust its phase when hearing a “fire” signal from a neighbouring robot. If there are more neighbours firing “adjustment signals” to a certain robot (i.e. we have a higher collective size $|R|$), it is also logical that the robot in question will update its phase more often. And so we can then see why e.g. all average harmonic synchronization times for $\alpha = 0.01$ seem to improve (i.e. gets reduced) for every increase in collective size $|R|$; the phase coupling constant $\alpha = 0.01$ might be weak, but with further and more frequent weak adjustments accumulated over the same time, the weak α is compensated for.

Hence, it seems like larger robot collectives do not require as large of a phase coupling constant α in order to synchronize to each other compared to that which smaller robot collectives require.

5.1.4 Comparing phase adjustment methods

Experiment setup:

So far in this chapter and section, we have only run experiments for pure phase synchronization, the ϕ problem, with Miroollo-Strogatz's method of phase adjustment, like described in 2.5.2.1. With this method of phase synchronization, robots are simply adjusting phases in an excitatory way; they only “push” other oscillators's phases further or higher when firing themselves, never “holding” or “dragging” them back.

Now, in order to investigate the validity of the claimed [] benefits of performing bi directional phase adjustments—both excitatory and inhibitory—like that of Nymoen's phase adjustment (see 3.2.2), an experiment comparing these two aforementioned phase adjustment methods thus follows.

Empirically speaking based on previous results in phase synchronization experiments, in particular the ones shown in Figure 5.1 and 5.2, arguably the best values so far for α (i.e. $\alpha = 0.8$) and t_{ref}^{dyn} (i.e. $t_{ref}^{dyn} = 0.1$) will be reused here for Miroollo-Strogatz's phase synchronization runs.

Since we have not yet performed any experiments for pure phase synchronization in the ϕ problem using Nymoen's slightly modified and bi directional phase adjustment function, the same α value of 0.8 will also be used for the phase synchronization runs here.

See the hyperparameter setup below in Table 5.4, and the results for the two phase adjustment methods given various robot collective sizes $|R|$ in 5.3. We here measure how long it takes various sizes of robot collectives to synchronize their phases to each other, given the two different phase adjustment methods.

α	β	t_{ref}^{dyn}	t_f	k	t_{max}
0.8	0	10%	80ms	8	5min

Table 5.4: Hyperparameter setup for comparative phase adjustment (Adj_ϕ) method ϕ synchronization experiment.

Experiment results:

We quickly see from the results in Figure 5.3 there seems to be a considerable difference in harmonic synchronization times between using Miroollo-Strogatz's and Nymoen's phase adjustment methods to synchronize phase in the phase (ϕ) synchronization problem. In Subfigure 5.3a, we see that for smaller musical robot collectives $|R|$, differences in synchronization time on average are not especially large. However, as collective size $|R|$ increases, we see the harmonic synchronization times for Nymoen's bi directional phase synchronization becomes considerably lower than that of Miroollo-Strogatz's mono-directional on average, variance—or at least the square root of it—accounted for.

On the other hand, when we look at the error rates in Subfigure 5.3b, we also see that Nymoen's phase adjustment method's error rates are considerably higher than those of Miroollo-Strogatz's error rates (being nearly negligible). Interestingly, this is only the case until collective size $|R| = 500$ though, after the point of which both phase adjustment methods's error rates are both equal to 0, and hence this objection towards Nymoen's phase adjustment method no longer holds. Hence, there seems to—at least up until $|R| = 500$ —exist a sort of tradeoff or high risk high reward scenario here, where one can in the phase (ϕ)

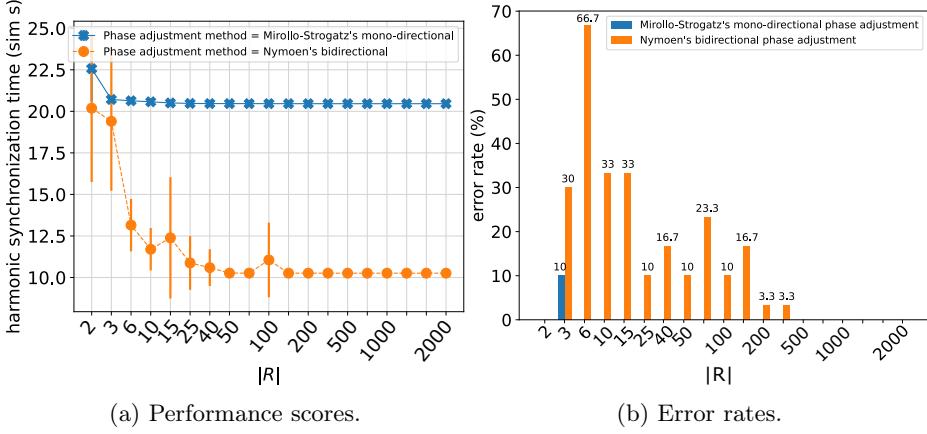


Figure 5.3: Performance scores 5.3a given in average harmonic synchronization times (sim s) with standard deviation. Corresponding error rates in 5.3b based on 30 individual runs per collective size $|R|$ for each of the two phase adjustment methods.

synchronization problem achieve harmonic synchrony faster by using Nymoen’s phase adjustment method while at the same time risking a higher chance of the synchronization failing; whereas Miroollo-Strogatz’s method of phase adjustment on average seems to be the safer, albeit slower, option.

Yet again, we also in these results notice a trend where smaller hyperparameter values lead to more unstable or worse performing synchronizations. Generally, although not strictly nor without exceptions, for both the harmonic synchronization times in Subfigure 5.3a and the error rates in 5.3b, the larger the hyperparameter $|R|$, the more stable or quick synchronization runs we performed in the Unity synchronization simulator.

Experiment analysis:

In this specific case, what this latter observation of seemingly increasing robustness given increasing collective sizes $|R|$ points to is perhaps that the system’s robustness and stability is increasing as collective size increases—being a strength classically advocated for in the multi agent systems and swarm systems literature [1]; as well as being found in e.g. robustness degrees in networks where you in the one case have a single point of failure versus many redundant paths and edges between your nodes [1].

As we see, the synchronization simulator in Unity is able to synchronize with a lot of agents, still without having broken the simulation yet. Comparatively, Nymoen et al. [22] only experimented with 6 oscillator nodes in their firefly system.

5.1.5 Increasing degree of self awareness

So far in the harmonic synchronization experiments, we have only and by default used the most restrictive and comprehensive *self awareness scopes*; all individual Dr. Squiggle robots have been globally connected, that is, and always hearing

every neighbour Dr. Squiggle's fire signals. Now we want to challenge this strong connectivity assumption.

Information and messages have the ability to travel through networks of many sorts, and the type of connectedness in these networks might influence how this occurs—and indeed the quality or accuracy of the travelled message; e.g. in a game of Chinese whispers versus a message sent through internet directly from the sender to the receiver. Both the speed as well as the overall quality of the message transmission can greatly be decided as a result of how connected the communicating nodes in a network are. Phenomena as these are to be experimented with here, as we will see how the connectedness between the musical Dr. Squiggle robots influence their collective and overall performance in synchronizing (harmonically).

It has previously, although more mathematically and not so much experimentally, been attempted to reduce the connectivity assumptions in pulse coupled oscillators [25] and to analyze the effects of it.

Here the hypothesis of whether increasing musical robots's degrees of self awareness will affect the synchronization performance or not, is investigated experimentally. Exactly what is meant by an increasing degree of self awareness specifically refers to the robots's *self awareness scope* [13]. This is tested in Unity for the more challenging $\phi\&\omega$ problem of harmonically synchronizing both phases and frequencies. Perhaps a larger self awareness scope, meaning more knowledge about the social environment, will lead to the robots having a better “overview” of the environment; hence leading to shorter simulation time (s) before reaching the goal state of harmonic synchrony. Or perhaps hearing more “fire” signals on average simply will be disturbing to the robots and hence disturb and slow down their entrainment towards harmonic synchrony. This experiments attempts to answer questions like these by for the three *self awareness scope* scenarios as introduced in 4.1.3.1 evaluating collective synchronization performance.

5.1.5.1 Self awareness scope tuning

In order to find out whether global connections in the pulse coupled oscillators is necessary to maintain performance in harmonic synchronization, in addition to finding some good values for a simulation breaking experiment where we will try to scale collective sizes $|R|$ as much as possible for robot collectives with varying self awareness scopes, varying *self awareness scopes* (as introduced in 4.1.3.1) are now experimented with.

k_s nearest neighbour SA scope experiment

Experiment setup:

Firstly, the first *self awareness scope* scenario (as in 4.1.3.1) is experimented with and tuned for. The number of neighbouring Dr. Squiggle robots each individual Dr. Squiggle robot will listen for fire signals to (k_s) will be scaled and increased from $k_s = 1$ up until $k_s = |R| - 1$ which is the maximum number of neighbours each Dr. Squiggle robot can listen to for fire signals.

In other words, the individual Dr. Squiggle robots's self awareness scope is increased incrementally in terms of k_s nearest neighbour SA scope and evaluated for performance in harmonic synchronization.

Hyperparameters are set in the simulator as shown in Table 5.5. Results can be seen in Figure 5.4.

α	β	t_{ref}^{dyn}	t_f	k	t_{max}
0.8	0	10%	80ms	8	5min

Table 5.5: Hyperparameter setup for k_s nearest neighbours SA scope ϕ synchronization experiment.

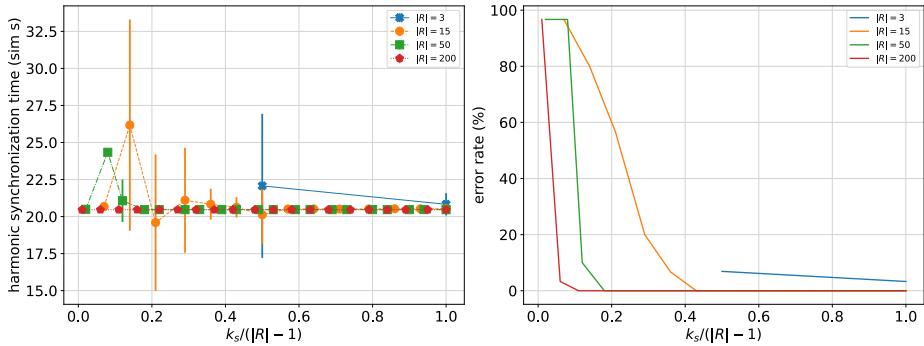


Figure 5.4: Performance times (left) given in average harmonic synchronization times (sim s) with standard deviation, given each nearest neighbour percentage $k_s/(|R|-1)$ for each $|R|$ and k_s value pair. Corresponding error rates (right) are based on 30 individual runs per collective size $|R|$ for each k_s nearest neighbour percentage evaluated for.

Experiment results:

As we can see on the x axis in Figure 5.4, we here plot harmonic synchronization scores per neighbour percentage $k_s/(|R|-1)$, which tells us the ratio of closest neighbours each robot listens to for fire signals out of all neighbours $|R|-1$ — or in other words the size of its self awareness scope, not given in absolute number of neighbours k_s , but in terms of the percentage of its neighbours.

d_s radial SA scope experiment

Experiment setup:

Results can be seen in Figure 5.5.

Experiment results:

As we see after $d_s = 65.6$, not much performance gain—if any—seems to be achieved by increasing the self awareness scope d_s further.

Experiment analysis:

Now, since we here tested with fixed d_s values instead of e.g. d_s -percentages like we did for k_s in experiment 5.4 above, it is clear that we will not get a increasing scaling of self awareness proportional to the varying collective sizes $|R|$. This is primarily due to how the spawning radius musical robots are instantiated within (cf. (4.2)) increases as collective sizes $|R|$ increases. This means

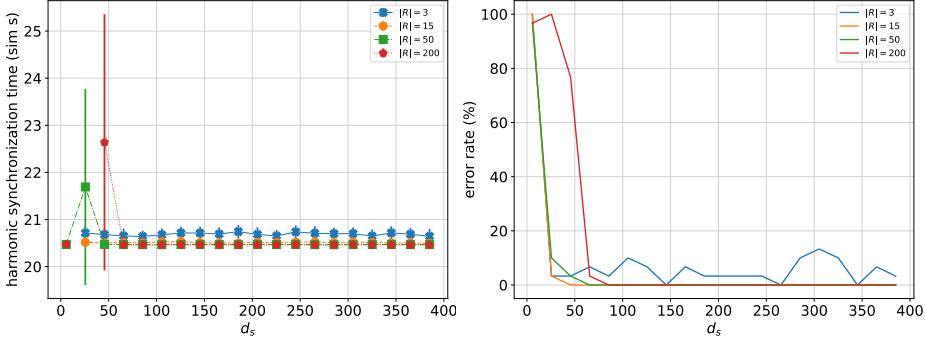


Figure 5.5: Performance times (left) given for various $|R|$ in average harmonic synchronization times (sim s) with standard deviation, given each SA scope radius d_s , spanning from the width of a robot to the maximum distance between robots in the largest robot collective evaluated for (i.e. $|R| = 200$). Corresponding error rates (right) are based on 30 individual runs per $|R|$ and d_s pair value evaluated for.

that even for the third d_s value (i.e. 45.6), collectives R with $|R| = 3$ having $radius_{spawn} = 13.65$ Unity physical units (cf. equation (4.2)) already practically have global self awareness scopes. For this reason and in hindsight, it might have been more fruitful to also here—instead of fixed d_s values per $|R|$ —test for neighbour percentages in radial self awareness scopes (i.e. $d_s/(|R| - 1)$).

5.1.5.2 Simulation breaking experiment

Experiment setup:

In order to test the limits of the developed synchronization simulator, we will here investigate the effects of scaling and exploding the musical robot collective size $|R|$ —also given some various self awareness scopes found worthwhile to further investigate in the previous experiments in 5.1.5.

The self awareness scope value of $k_s = 0.1$ was chosen in order to have a sense of measurement and comparison, in order to see that in fact robot collectives might still manage to synchronize even though they listen to only one neighbouring robot's fire signals for self adjustment; but probably with way less performance comparatively to larger self awareness scopes. Furthermore, the value of $d_s = 75$ is chosen to be further experimented with, as not much improvement was seen after this point for robot collectives in the relevant results in Figure 5.5. And finally, the maximum size of self awareness scopes there is, the global self awareness scopes were set in the individual robots in the collective.

Additional experiment hyperparameters apart from $Adj_\phi = \text{Miroollo \& Strogatz}$ and $Adj_\omega = \text{None}$, as mentioned in the start of Section 5.1, can be seen in Table 5.6. Experiment results after running simulations in Unity can thus be seen in Figure ??.

α	β	t_{ref}^{dyn}	t_f	k	t_{max}
0.8	0	10%	80ms	8	5min

Table 5.6: Hyperparameter setup for collective size explosion ϕ synchronization experiment.

Experiment results:

FYLL INN!

Experiment analysis:

5.1.5.3 Heterogenous self awareness scopes experiment

Experiment setup:

In order to test whether actually globally connected oscillators are necessary to maintain harmonic synchronization performance, or whether more relaxed and less constrained connections in the oscillators still lead to similar performance in harmonic synchronization performance—heterogenous robot collectives with varying degrees of self awareness scopes will be explored here.

First, relatively un-aware robot collectives (i.e. where musical robots have small or narrow self awareness scopes) are given the task of harmonically synchronizing to each other. Then, a more neighbour-aware robot collective (with some robots having larger self awareness scopes) is to harmonically synchronize; and so on for more and more self aware robot collectives.

We will see whether the intuition that more knowledge about ones neighbourhood or environment will lead to increased performance, or whether performance is still maintained despite less awareness in robots of their neighbourhood.

The nearest neighbour SA scope value of $k_s/(|R| - 1) = 0.4$ will be used in this experiment for the smallest or minimal self awareness scope—referred to as the *minimal SA scope*. This value of 0.4 (corresponding to 40% of each robot's closest neighbour robots) is chosen due to its good results in Figure 5.4, as performance is barely increased after that point. Consequently, the collective size $|R| = 15$ is selected for this experiment again due to this value's performance results in Figure 5.4; in Figure 5.4, collective size $|R| = 15$ yielded the most interesting and varying behaviour, and compared to larger collective sizes one can with $|R| = 15$ not simply blame good performance on the “strength of numbers”, as larger collective sizes generally seem to synchronize more steadily. The radial self awareness scope value $d_s = 65.6$ is chosen due to the lack of further improvement in performance scores in the relevant experiment shown results for in Figure 5.5. However; given the lesson learned that it might be more fruitful to test for neighbour percentages when it comes to self awareness scopes, the corresponding percentage for $|R| = 200$ and $d_s = 65.6$ in the aforementioned experiment (cf. Figure 5.5), is chosen for the *radial self awareness scope* in this experiment. Hence, we introduce the radial SA scope percentage as follows:

$$d_s^{\%} := d_s / (2 \cdot radius_{spawn}), \quad (5.1)$$

where d_s is the radius (physical Unity units) within which the robot in question will listen to neighbours's fire signals for adjusting itself, and the robot collective instantiation radius $radius_{spawn}$ is defined in equation (4.2). By using the found values of $d_s = 65.6$ and $|R| = 200 \implies radius_{spawn} = 194$, we then acquire our desired $d_s^{\%} = 65.6 / (2 \cdot 194) = 0.17 = 17\%$.

In other words, if a musical robot has a radial SA scope of $d_s\% = 17\%$, it implies that this robot will listen to (for “fire” signals) from all neighbouring robots with a radius of less than 17% times the diameter of the spawning area (i.e. $2 \cdot \text{radius}_{\text{spawn}}$).

And finally, robots will increasingly also have *global self awareness scopes*, as these are the most *neighbour aware* robot collectives.

Hence, and since we now will give heterogenous robot collectives the task of synchronizing harmonically to each other, robot collectives’s individual robot self awareness scopes will now as a whole consist of ratios on the form:

$$(\text{minimal}, \text{radial}, \text{global}) = (m, r, g) \quad (5.2)$$

where the *minimal*, *radial*, and *global* self awareness scopes are as defined above. E.g., the least self aware robot collective imaginable would in this case be one of SA scopes $(1.0, 0.0, 0.0)$ (i.e. $|R| * 1.0$ robots would have the *minimal* SA scope, and $|R| * 0.0 = 0$ robots would have *radial* or *global* SA scopes), whereas the most imaginable self aware robot collective would have self awareness scopes $(0.0, 0.0, 1.0)$. These two will serve as the boundaries within which we want to test for various percentages of self awareness scopes, given the three percentages m , r , and g .

The other hyperparameter values for this heterogenous SA scopes experiment is given in Table 5.7. The synchronization results obtained by running experiments in the Unity simulator is shown in Figure ??.

α	β	t_{ref}^{dyn}	t_f	k	t_{max}
0.8	0	10%	80ms	8	5min

Table 5.7: Hyperparameter setup for heterogenous SA scopes ϕ synchronization experiment.

Experiment results:
[FYLL INN!](#)

Experiment analysis:

5.2 Phase and frequency synchronization

This is the section for the experiments attempting to synchronize for the second and harder problem of synchronizing both phases ϕ_i , as well as frequencies ω_i , for all agents i . These are then experiments where all musical robots originally have unequal and ever-changing phases and frequencies, adjusting both phases and frequencies in order to entrain to synchronize their phases and frequencies until reaching harmonic synchrony. Or in other words, all experiments in this phase and frequency synchronization section have the two hyperparameters shown in Table 5.8 in addition to the individual experiment hyperparameters listed; simply reflecting that in this problem we have to use both a phase adjustment (Adj_ϕ) function as well as a frequency adjustment (Adj_ω) function.

Adj_{ϕ}	Adj_{ω}
Nymoen	Nymoen

Table 5.8: Hyperparameter setup for phase and frequency (ϕ & ω) synchronization experiments.

5.2.1 Initial validation

To see that synchronization is actually happening like intended, some relevant quantities and simulation values are recorded for a phase synchronization simulation run in Unity and displayed in ??.

FYLL INN.

5.2.2 Reproducing baseline results

Again, we want to here see whether we get more or less the same results in the Unity simulator as Nymoen get in their firefly oscillator system.

Hence, we here present attempts made in the novel Unity synchronization simulator at recreating Nymoen's first results with their novel frequency synchronization method, which is utilizing, amongst other aspects, self awareness [22].

5.2.2.1 Ordering by phase couplings experiment

Experiment setup:

Given that Nymoen do not mention their β value in their frequency synchronization experiment where they order for different phase coupling values α , an *empirically decent* β value of 0.4 is chosen for this experiment. What *empirically decent* refers to in this case are Nymoen's findings in the results of their last experiment [22] where synchronization times for various β values were evaluated; deeming $\beta = 0.4$ to be a good value, with no further improvement in synchronization performance when β is increased further.

Here, Nymoen's self aware frequency adjustment method, implemented in our novel synchrony simulator in Unity, is experimented with for varying phase coupling constants α . This time not only oscillator phases are initially unsynchronized; oscillator frequencies are also unsynchronized to begin with. Hence, we here try to synchronize our musical robots in the phase (ϕ) *and* frequency (ω) synchronization problem, using Nymoen's phase adjustment method to synchronize phases, as well as Nymoen's frequency adjustment method to synchronize frequencies. See set up hyperparameters in Table 5.9. See the results in Figure 5.6.

$ R $	β	ω_{min}^{init}	ω_{max}^{init}	m	t_{ref}	t_f	k	t_{max}
6	0.4	0.5Hz	4Hz	5	50ms	80ms	8	5min

Table 5.9: Hyperparameter setup for first ϕ & ω synchronization baseline reproduction experiment.

Experiment results:

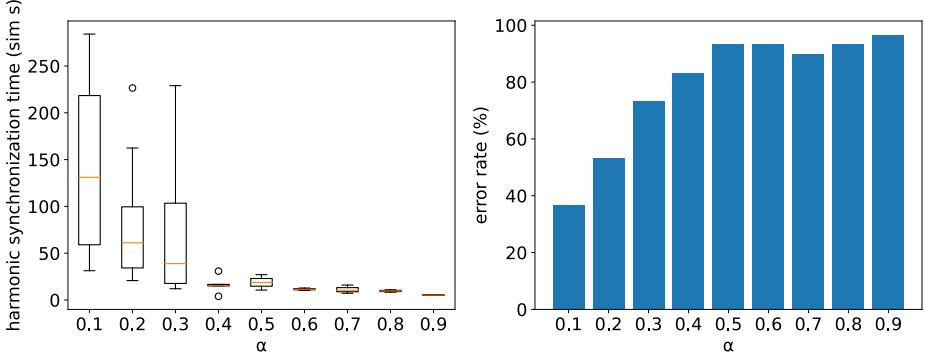


Figure 5.6: Harmonic synchronization times (sim s) for 6 robots with both initially unsynchronized phases *and* frequencies, for varying phase coupling constants α , reaching harmonic synchrony—but also often failing to. 30 simulation runs per α are reported.

5.2.2.2 Hopefully stabler ordering by phase couplings experiment

Experiment setup:

After seeing how poorly the musical robot collectives managed to achieve harmonic synchrony during previous experiment 5.2.2.1, we change the hyperparameters slightly in the hopes of achieving harmonic synchrony more frequently, or in other words more stable synchronization simulation runs. The reason is mostly that it would be beneficial to actually see whether we observe similar patterns as in Nymoen’s results—something which becomes impossible when robot collectives nearly never achieve harmonic synchrony.

The frequency coupling constant of $\beta = 0.6$ could potentially be another supposed good β value, at least solely judging by Nymoen’s results, and hence we select this value for this hopefully stable phase and frequency synchronization experiment.

By initial trial and error, it also became apparent that phase *and* frequency synchronization was not always stable—i.e. the robot collective not managing to achieve harmonic synchrony within the maximum time limit—for collective sizes of 6 or more. It did however become apparent that phase *and* frequency synchronization *was* stable for smaller musical robot collective sizes, like 2 or 3.

Hence, a similar experiment to as in 5.2.2.1 is set up in the Unity synchrony simulator, and is run similarly as before. So still, the phase (ϕ) *and* frequency (ω) synchronization problem is to be experimented for, given different phase coupling constants α , with Nymoen’s both phase and frequency adjustment methods—only this time with $\beta = 0.6$ and *collsize* = 3 instead. See set up hyperparameters in Table 5.10. See results in Figure 5.7.

$ R $	β	ω_{min}^{init}	ω_{max}^{init}	m	t_{ref}	t_f	k	t_{max}
3	0.6	0.5Hz	4Hz	5	50ms	80ms	8	5min

Table 5.10: Hyperparameter setup for ϕ & ω synchronization baseline reproduction and stabilization experiment.

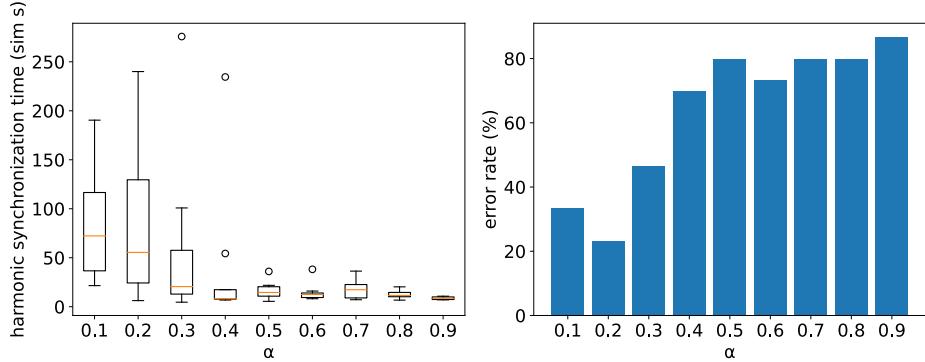


Figure 5.7: Harmonic synchronization times (sim s) for 3 robots with both initially unsynchronized phases *and* frequencies for varying phase coupling constants α . 30 simulation runs per α are reported.

Experiment results:

5.2.2.3 Ordering by frequency couplings experiment

Experiment setup:

Lastly, for the final Baseline reproducing experiment, it was wanted to see whether results within the phase and frequency ($\phi \& \omega$) synchronization problem, when one varied the frequency coupling constant β , also was more or less in the same ballpark or not, for the same reasons as mentioned in the first Baseline reproducing experiment in Section 5.1.

Since we now do not synchronize in Unity for varying phase coupling constants α , we now fix α and rather test how the individual musical robots's various frequency coupling constants β affect the performance of the musical robot collective.

Again, Nymoen does not specify exactly the phase coupling constant α they use in their latter experiment when testing their firefly-inspired synchronization system for various β values. Hence, the now fixed phase coupling constant α is here selected by reusing a reasonably good α value, based on the similar $\phi \& \omega$ synchronization experiments presented in Figure 5.6 and 5.7. In these two experiments in particular, $\alpha = 2$ yielded both harmonic synchronization times among the best synchronization times, given the error scores; and furthermore, one of the lowest error rates. This then gives us a fixed $\alpha = 0.2$.

See the set up hyperparameters for the experiment in Table 5.11, and the corresponding results in Figure 5.8.

$ R $	α	ω_{min}^{init}	ω_{max}^{init}	m	t_{ref}	t_f	k	t_{max}
6	0.2	0.5Hz	4Hz	5	50ms	80ms	8	5min

Table 5.11: Hyperparameter setup for second ϕ & ω synchronization baseline reproduction experiment.

Experiment results:

Even though we do not see exactly the same harmonic synchronization times

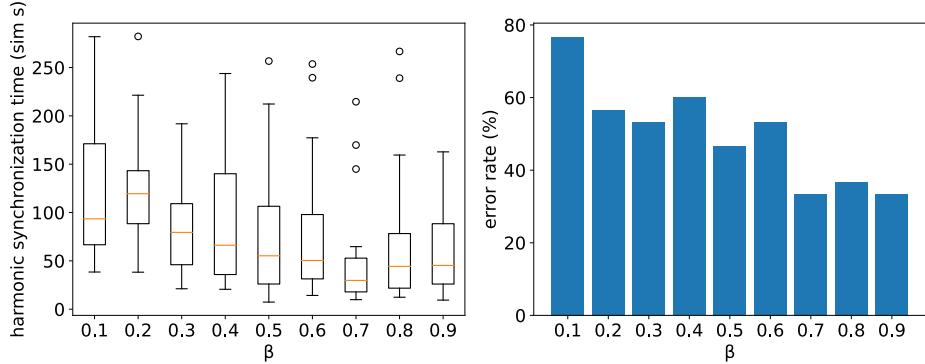


Figure 5.8: Synchronization times (s) for 6 robots with both initially random and unsynchronized phases, and frequencies, for varying frequency coupling constants β . 30 simulation runs per β are reported.

as we see in Nymoen’s results, and worse at that, we do in fact see a similar pattern in that also here, synchronization times and error scores seem to improve the larger frequency coupling constant β we have.

Experiment analysis:

Also, from the results in Figure 5.8, it becomes apparent that $\beta = 0.7$ might be the best choice to continue using in our Unity synchrony simulator—at least when it comes to collective sizes $|R| = 6$.

5.2.3 Increasing degree of self awareness

Motivations and intentions are exactly the same as in 5.1.5, although now it is time to explore the effects of various degrees and scopes of self awareness in the second and more challenging phase and frequency ($\phi\&\omega$) synchronization problem.

5.2.3.1 Self awareness scope tuning

Similarly to in the experiments in 5.1.5.1 we here want find out whether global connections in the pulse coupled oscillators really is necessary to maintain performance in harmonic synchronization, but for this time in the second and more complex ($\phi\&\omega$) synchronization problem.

k_s nearest neighbour SA scope experiment

Experiment setup:

The first *self awareness scope* scenario (as in 4.1.3.1) is now again experimented with and tuned—only now for the phase and frequency ($\phi\&\omega$) synchronization problem. The number of neighbouring Dr. Squiggle robots each individual Dr. Squiggle robot will listen for fire signals to (k_s) will be scaled, for each $|R|$, and increased from $k_s = 1$ up until $k_s = |R| - 1$ which is the maximum number of neighbours each Dr. Squiggle robot can listen to for fire signals.

In other words, the individual Dr. Squiggle robots's self awareness scope is increased incrementally in terms of k_s nearest neighbour SA scope and evaluated for performance in harmonic synchronization.

Hyperparameters are set in the simulator as shown in Table 5.12. These are empirically found to be the best hyperparameters in previous experiment results in Section 5.2 where we are attempting synchronization in the $\phi\&\omega$ problem.

Results can be seen in Figure 5.9.

α	β	t_{ref}^{dyn}	t_f	k	t_{max}
0.2	0.7	10%	80ms	8	5min

Table 5.12: Hyperparameter setup for k_s nearest neighbours SA scope ϕ & ω synchronization experiment.

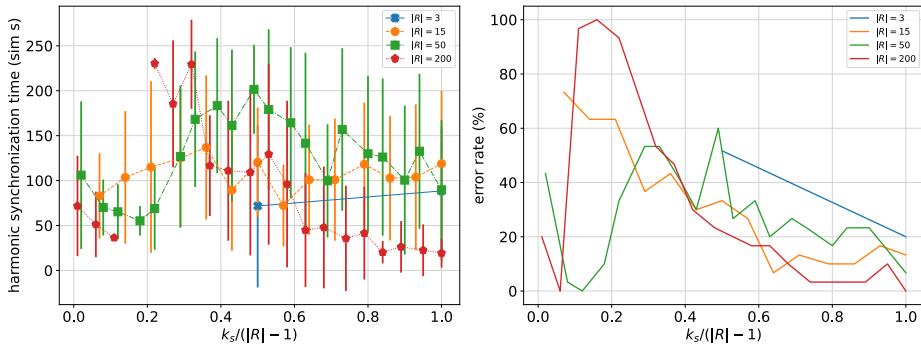


Figure 5.9: Performance times (left) given in average harmonic synchronization times (sim s) with standard deviation, given each nearest neighbour percentage $k_s/(|R|-1)$ for each $|R|$ and k_s value pair. Corresponding error rates (right) are based on 30 individual runs per collective size $|R|$ for each k_s nearest neighbour percentage evaluated for.

Experiment results:

Also here, one can see on the x axis in 5.4 that harmonic synchronization scores are plotted per neighbour percentage $k_s/(|R|-1)$, which tells us the ratio of closest neighbours each robot listens to for fire signals out of all neighbours $|R|-1$ — or in other words the size of its self awareness scope, not given in absolute number of neighbours k_s , but in terms of the percentage of its neighbours. Just like in the similar experiment for the phase (ϕ) synchronization problem in 5.1.

Experiment analysis:

As is visible in the synchronization times and error rates in Figure 5.9, often times musical robot collectives still manage to achieve harmonic time synchronization despite listening to smaller percentages of their neighbour robots. From around after musical robots r_i in the collectives R have 20% of their neighbouring robots in their self awareness scope (i.e. $k_s/(|R|-1) = 0.2$), there seems to be a trend we might expect; namely, synchronization times increase and error rates decrease as we increase the self awareness scope of the musical robots.

As individual musical robots listen to more and more of their neighbouring Dr. Squiggle robots for synchronizing themselves upon their neighbour’s “fire” signal, the faster the musical robot collective as a whole manage to harmonically synchronize—and with less synchronization fails.

However, there actually also seems to be a region, when and before robots in the robot collectives have $\approx 10\%$ of their neighbours in their SA scopes, where harmonic synchronization performance is good both in terms of synchronization time as well as low error rates.

Not much improvement in harmonic synchronization time or error rate is seen beyond the point where individual musical robots r_i have 80% of their neighbours in their self awareness scopes, and so perhaps it might be sufficient in musical robot collectives that each robot listens and is connected to 80% of its closest neighbours, relaxing the globally connectivity assumption.

d_s radial SA scope experiment

Experiment setup:

Results can be seen in Figure 5.10.

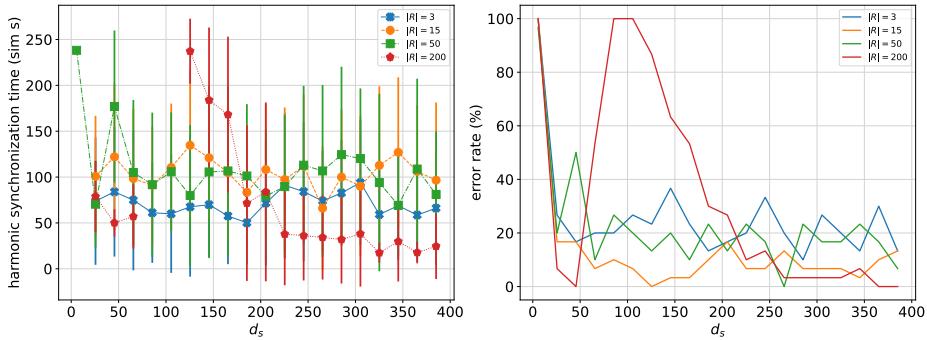


Figure 5.10: Performance times (left) given for various $|R|$ in average harmonic synchronization times (sim s) with standard deviation, given each SA scope radius d_s , spanning from the width of a robot to the maximum distance between robots in the largest robot collective evaluated for (i.e. $|R| = 200$). Corresponding error rates (right) are based on 30 individual runs per $|R|$ and d_s pair value evaluated for.

Experiment results:

Experiment analysis:

As was alluded to before, as well as heavily experienced while developing and performing phase *and frequency* synchronization in Unity, can also be seen in Figure 5.10; namely, the synchronization process is much less stable and predictable, as we see greatly overlapping standard deviations in the harmonic synchronization time plot.

However, despite the instability, harmonic time synchronization in the musical robot collective is achieved, although as we can see in the error rate plot

with much higher errors on average than with pure phase synchronization, as e.g. in Figure 5.5.

Also note that due to how the musical robots are instantiated in the Unity scene, for lower collective sizes $|R|$, radial SA scopes effectively become global SA scopes; e.g. if a robot collective has a diameter across the instantiation / spawning area of only 13.65 Unity physical units (which is the case for robot collectives of $|R| = 3$, cf. (4.2)), and the radial self awareness scope $d_s = 50$.

Chapter 6

Conclusions

BESKR.: [Where I shall follow-up on my *research questions* by a discussion of to which degree—and in what ways—the thesis-/project-work has answered them].

GJØR: [Se på (for kapittel-inspirasjon):

- Tønnes . MSc-thesis . Discussion-Ch.
- Jim . 'how to write a master thesis.pdf' . 'Conclusions'.

].

6.1 General discussion

As we can see from the results in the previous chapter, musical robot collectives achieve harmonic synchronization with various performance dependant on various robot collective or individual hyperparameter values.

We see e.g. the recurring trend that the larger the robot collectives R we have (i.e. the larger $|R|$ we have), the more stable the robot collective seems to manage synchronizing to each other—compared to with lower collective sizes $|R|$.

Perhaps not so surprisingly, performance both in terms of harmonic synchronization time as well as low error rates also tends to improve as we endow individual musical robots with larger self awareness scopes; the more of the closest neighbourhood each robot is self aware of, the higher the performance tends to be. However, there were also cases where not much improvement was seen in harmonic synchronization performance past certain points in self awareness scopes—even before robots having all their neighbours in their scope (i.e. global SA scope).

6.1.1 Conclusion

To summarize, time synchronization (harmonic at that) in musical robots, modelled as pulse coupled oscillators, has been achieved in the newly developed

synchronization simulator, utilizing some already existing methods for synchronizing both oscillator phases and oscillator frequencies.

Although individual hyperparameters plays a role in terms of harmonic synchronization performance to a certain extent, what is often more decisive for in a stable manner achieving harmonic synchrony in an oscillator collective is the collective size $|R|$.

Furthermore, musical robots need not to be globally connected in order to achieve harmonic synchronization often times, and one can even relax the connectivity whilst still achieving as good synchronization performance as globally connected robots.

However, there still is an overarching trend of increasing harmonic synchronization performance the larger the self awareness scope of the individual musical robots are. Such increased degree of *neighbour awareness* leading to higher performance in a multi agent system was also found by J. Cao et al. [2] as they managed to increase cooperation in a multi-agent reinforcement learning (MARL) scenario in the challenging videogame StarCraft II. This they achieved by increasing each agent's local awareness of its neighbour. To put this finding into a broader context, and a bit outside of what is studied here, this trend of increased cooperation and overall performance through *increased neighbour self awareness* might be a good argument for why one should always be open to input—both in life, organizations, and businesses—as well as widen ones horizon and to be open minded. As Jordan Peterson's ninth rule in his book *12 Rules for Life* reads: “Assume that the person you are listening to might know something you don't.”

6.1.2 Ethical reflections

One might get frightened by hearing robots or machines have become self aware. This can remind listeners of concepts like superintelligence or general artificial intelligence, which can have scary and unknown consequences if ever achieved.

My personal view is that the work presented in this thesis probably holds potential for computing systems and robots to get more well defined and concise self awareness specifications, and hence might become more self aware as a consequence of such clarity. Furthermore, frameworks for computing systems inspired by psychology in terms of self awareness might potentially get us closer to achieving something resembling general artificial intelligence or a superintelligence. However, I do not fear of this happening in the near future or even the far future; I would like to echo what Andrew Ng, being one of the leading figures within the AI community, has been quoted to say regarding the possibility of an AI superintelligence: “Worrying about AI evil superintelligence today is like worrying about overpopulation on the planet Mars. We haven't even landed on the planet yet!” I, like others, even think there is a safety argument to be had about self awareness having a larger focus in computational systems like robots, being that explainability (being a hot topic at the moment) can be further increased by considering computational self awareness in a computing system. However, I am very in favor of careful considerations and regulatory oversight from all parts of society in order for humans to ensure control and desired outcomes when it comes to both AI, and more specifically to topics like the one explored in this thesis being computational self awareness.

6.2 Further work

Physical domain: Going from simulation to the real physical world, using M. J. Krzyzaniak and RITMO’s actual musical robots, the Dr. Squiggles (cf. 2.2.1.2). This would involve implementing or translating the software written for the simulated Dr. Squiggles in the Unity simulator environment into the physical hardware and interfaces available in the real Dr. Squiggle robots. Using an audio interface with the Dr. Squiggles robots in order to transmit or play audible “fire” signals would perhaps correspond to *self awareness scope 2* as defined in this thesis, whereas *self awareness scope 3* could possibly correspond in reality to communication technologies where messages are heard globally relatively instantaneously after it is transmitted, such as e.g. wireless local area networks (e.g. Wi-Fi) or—even faster—Ethernet internet cables.

Introduction of noise and epistemological challenges: In this music system, pulses (or firings) were designed in such a way that the other musical nodes would easily hear the distinct signals. We could have easily imagined a scenario where this would not be given; hence calling for more sophisticated methods for detecting and separating between actually separate (spatially) audio-sources. Epistemological explorations, related to e.g. Computational Auditory Scene Analysis (as mentioned in [19]) would then be of interest. Some simple ways to introduce noise in the developed synchronization simulator would be to omit or skip, with a probability, notifying a robot by sending a “fire” or adjustment signal to it for self adjustment. In this way, it would be like robots randomly not managing to listen for all the “fire” signals being transmitted into the environment. One could then do a study on how well different musical robot collectives, in terms of e.g. self awareness scopes, would handle this additional challenge.

Automatic and online (harmonic) fire sound generation: As alluded to in 4.1.3.2, one possible avenue to explore in order to find harmonious chords or tones—when played together—in a more automatic approach, live and online during simulation, could be found through musical machine learning models predicting fitting and accompanying chords and harmonies (given a tone) in real-time, like the ones tested thoroughly in B. Wallace’s interactive music system PSCA [30]. This was however not attempted in this thesis, as this was beyond the scope of the project.

Hindrances throughout the thesis work: The main reason why these aforementioned ideas were not possible to explore throughout this thesis work in the end, was due to an (to me at least) unseen and—for months—unobservable Unity peculiarity (perhaps only to a Unity novice like me) which in turn caused visible bugs in two mechanisms—namely the harmonic synchrony detection mechanism as well as my implementation of K. Nymoen et al.’s frequency adjustment method. So given that I only saw problems in the harmonic synchrony detection and frequency synchronization, I ended up spending considerable time throughout the thesis period trying to debug these two mechanisms, without success—or if any then with more bugs popping up. The way out of this extremely confusing and frustrating situation reared its head eventually as my attention was brought to the synchrony simulator’s determinism. Little did I

know that once I made my synchrony simulator deterministic, the problems I had been debugging for months were solved immediately. And after discussing with supervisors, and a much more experienced Unity user (Frank Veenstra), the root (and to me unseen) problem, and the reasons why it lead to those visible problems in my harmonic synchrony detection and frequency synchronization, suddenly made total sense. And so now this resulting knowledge about indeterminism in Unity, learnt the hard way, will most likely be shared through a ROBIN wiki page or the likes. Hopefully, this will prevent future UiO robotics master's students from stumbling onto the same problems I did, and from ending up in similar rabbit holes as I ended up in.

Bibliography

- [1] Jens B Asendorpf, Veronique Warkentin, and Pierre-Marie Baudonnier. “Self-awareness and other-awareness. II: Mirror self-recognition, social contingency awareness, and synchronic imitation.” In: *Developmental Psychology* 32.2 (1996), p. 313.
- [2] Jiahao Cao et al. “LINDA: Multi-Agent Local Information Decomposition for Awareness of Teammates”. In: *arXiv preprint arXiv:2109.12508* (2021).
- [3] Arjun Chandra et al. “Enabling participants to play rhythmic solos within a group via auctions”. In: *9th International Symposium on Computer Music Modelling and Retrieval (CMMR 2012)*, Queen Mary University of London (2012).
- [4] Arjun Chandra et al. “Reference Architecture for Self-aware and Self-expressive Computing Systems”. In: *Self-aware Computing Systems: An Engineering Approach*. Ed. by Peter R. Lewis et al. Natural Computing Series. Cham: Springer International Publishing, 2016. Chap. 4, pp. 37–49. ISBN: 978-3-319-39675-0. DOI: [10.1007/978-3-319-39675-0_4](https://doi.org/10.1007/978-3-319-39675-0_4).
- [5] Louise A. Dennis and Michael Fisher. “Verifiable self-aware agent-based autonomous systems”. In: *Proceedings of the IEEE* 108.7 (2020). Publisher: IEEE, pp. 1011–1026.
- [6] Marco Dorigo, Mauro Birattari, and Thomas Stutzle. “Ant colony optimization”. In: *IEEE computational intelligence magazine* 1.4 (2006). Publisher: IEEE, pp. 28–39.
- [7] Shelley Duval and Robert A Wicklund. “A theory of objective self awareness.” In: (1972).
- [8] Daniel Goldman, Haldun Komsuoglu, and Daniel Koditschek. “March of the sandbots”. In: *IEEE Spectrum* 46.4 (2009). Publisher: IEEE, pp. 30–35.
- [9] Pentti OA Haikonen. “Reflections of Consciousness: The Mirror Test.” In: *AAAI Fall Symposium: AI and Consciousness*. 2007, pp. 67–71.
- [10] Auke Jan Ijspeert. “Central pattern generators for locomotion control in animals and robots: a review”. In: *Neural networks* 21.4 (2008). Publisher: Elsevier, pp. 642–653.
- [11] Keiji Konishi and Hideki Kokame. “Synchronization of pulse-coupled oscillators with a refractory period and frequency distribution for a wireless sensor network”. In: *Chaos: an interdisciplinary journal of nonlinear science* 18.3 (2008). Publisher: American Institute of Physics, p. 033132.

- [12] Samuel Kounev et al. “The Notion of Self-aware Computing”. In: *Self-Aware Computing Systems*. Ed. by Samuel Kounev et al. Cham: Springer International Publishing, 2017. Chap. 1, pp. 3–16. ISBN: 978-3-319-47474-8. DOI: 10.1007/978-3-319-47474-8_1.
- [13] Peter Lewis et al. “Towards a Framework for the Levels and Aspects of Self-aware Computing Systems”. In: *Self-Aware Computing Systems*. Ed. by Samuel Kounev et al. Cham: Springer International Publishing, 2017. Chap. 3, pp. 51–85. ISBN: 978-3-319-47474-8. DOI: 10.1007/978-3-319-47474-8_3.
- [14] Peter R. Lewis, Arjun Chandra, and Kyrre Glette. “Self-awareness and Self-expression: Inspiration from Psychology”. In: *Self-aware Computing Systems: An Engineering Approach*. Ed. by Peter R. Lewis et al. Natural Computing Series. Cham: Springer International Publishing, 2016. Chap. 2, pp. 9–21. ISBN: 978-3-319-39675-0. DOI: 10.1007/978-3-319-39675-0_2.
- [15] Renato E. Mirollo and Steven H. Strogatz. “Synchronization of pulse-coupled biological oscillators”. In: *SIAM Journal on Applied Mathematics* 50.6 (1990). Publisher: SIAM, pp. 1645–1662.
- [16] Melanie Mitchell. “Self-awareness and control in decentralized systems.” In: *AAAI Spring Symposium: Metacognition in Computation*. 2005, pp. 80–85.
- [17] Francesco Mondada et al. “SWARM-BOT: A new distributed robotic concept”. In: *Autonomous robots* 17.2 (2004). Publisher: Springer, pp. 193–221.
- [18] Alain Morin. “Levels of consciousness and self-awareness: A comparison and integration of various neurocognitive views”. In: *Consciousness and cognition* 15.2 (2006), pp. 358–371.
- [19] Kazuhiro Nakadai and Hiroshi G. Okuno. “Robot audition and computational auditory scene analysis”. In: *Advanced Intelligent Systems* 2.9 (2020). Publisher: Wiley Online Library, p. 2000050.
- [20] Z Néda et al. “SelfOrganizing processes: The sound of many hands clapping”. In: *Nature* 403 (2000), p. 6772.
- [21] Ulric Neisser. “The Roots of Self-Knowledge: Perceiving Self, It, and Thou a”. In: *Annals of the New York Academy of Sciences* 818.1 (1997), pp. 19–33.
- [22] Kristian Nymoen et al. “Decentralized harmonic synchronization in mobile music systems”. In: *Awareness Science and Technology (iCAST), 2014 IEEE 6th International Conference on*. IEEE, 2014, pp. 1–6.
- [23] Yvan R. Petillot et al. “Underwater robots: From remotely operated vehicles to intervention-autonomous underwater vehicles”. In: *IEEE Robotics & Automation Magazine* 26.2 (2019). Publisher: IEEE, pp. 94–101.
- [24] Anton V. Proskurnikov and Ming Cao. “Synchronization of pulse-coupled oscillators and clocks under minimal connectivity assumptions”. In: *IEEE Transactions on Automatic Control* 62.11 (2016). Publisher: IEEE, pp. 5873–5879.

- [25] Anton V. Proskurnikov and Ming Cao. “Synchronization of pulse-coupled oscillators and clocks under minimal connectivity assumptions”. In: *IEEE Transactions on Automatic Control* 62.11 (2016). Publisher: IEEE, pp. 5873–5879.
- [26] Bernhard Rinner et al. “Self-aware and self-expressive camera networks”. In: *Computer* 48.7 (2015). Publisher: IEEE, pp. 21–28.
- [27] Thomas Schmickl et al. “CoCoRo—The Self-Aware Underwater Swarm”. In: *2011 Fifth IEEE Conference on Self-Adaptive and Self-Organizing Systems Workshops*. IEEE, 2011, pp. 120–126.
- [28] Fikret Sivrikaya and Bülent Yener. “Time synchronization in sensor networks: a survey”. In: *IEEE network* 18.4 (2004). Publisher: IEEE, pp. 45–50.
- [29] The Macat Team. *An Analysis of William James’s The Principles of Psychology*. London: Macat Library, July 15, 2017. 95 pp. ISBN: 978-1-912282-49-4. DOI: [10.4324/9781912282494](https://doi.org/10.4324/9781912282494).
- [30] Benedikte Wallace. “Predictive songwriting with concatenative accompaniment”. Master’s Thesis. 2018.